

Lectures

Introduction

L0 and L1

Program Parallelization

Decomposition: Decompose a sequential algorithm into tasks (programmer)

- Granularity of tasks are important
- Tasks have dependencies (data or control) between each other which defines the execution order

Scheduling: Assign tasks to processes (programmer / compiler)

Mapping - Map processes to cores (OS)

Von Neumann Computation Model instruction and data are stored in memory, and processors computes.

Memory Wall disparity between memory speed and processor speed (≤ 1 ns VS ≥ 100 ns)

Processing unit refers to a core that can execute a kernel thread

Interconnect busses from different components in the machine

Node Machine in a distributed system

Why Parallel

Primary Reasons

- Overcome limits of serial computing
- Solve larger problems
- Save (wall-clock) time

Other Reasons

- Take advantage of non-local resources
- Cost/energy saving - use multiple cheaper computing resources
- Overcome memory constraints

Computational Model Attributes


- Operation mechanism** Primitive units of computation or basic actions of the computer on a specific Architecture

- Data Mechanism** How we access and store data in address space
- Control Mechanism** How primitive units of computation are scheduled
- Communication Mechanism** Modes and patterns of exchanging information between parallel tasks (e.g message passing, shared memory)
- Synchronization Mechanism** ensures to ensure needed information arrives at the right time

Dependencies and Coordination

- Dependencies among tasks impose constraints on scheduling
- Memory organizations: Shared-memory (threads), distributed-memory (processes)
- Coordination (synchronisation) imposes additional overheads

Two algorithms



- Core 0 is active throughout the execution
- Some cores are idle
- This is a lot better than having all cores idle while the master core is executing

Parallel Performance

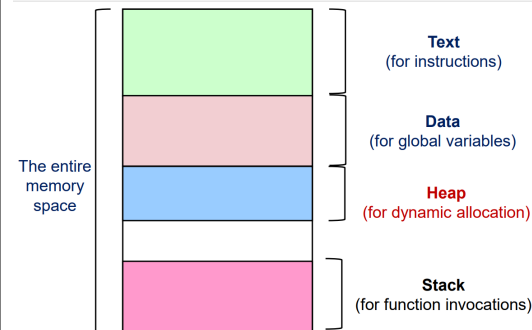
- Execution time Vs Throughput
- Parallel execution time = computation time + parallelization overheads
- Overheads: Distribution of work(tasks) to porocesses, information exchange, synchronisation, idle time, etc

Background on Parallelism

L2: Processes and Threads

Process

- Identified by PID
- Program counter, global data (open files, network connections), stack or heap, current values of the registers (GPRs and Special)
- These information are abstracted in the PCB, and each proecss can be viewed as having exclusive access to tis address space
- Explicit communication is needed
- Disadvantage**
 - High overhead of system calls
 - Potential re-allocation of data-structures
 - Communication goes through OS (system calls) and context switch is costly



Multi tasking

- Overhead: Context switching (PCB change) is needed and states of suspended process must be saved
- Time slicing: Pseudo-parallelism
- Child processes can use parent's data

Inter-process communication (IPC)

- Shared memory: need to protect access with locks
- Message passing: Blocking, unblocking, Synchronous, unsynchronous

Exceptions

- Executing a **machine level instruction** can cause exception
- For example: Overflow, Underflow, Division by Zero, Illegal memory address, Mis-aligned memory access

Synchronous

- Occur due to program execution
- Have to execute an **exception handler**

Interrupts

- External events can interrupt the execution of a program
- Usually hardware related: Timer, Mouse Movement, Keyboard Pressed etc

Asynchronous

- Occur **independently** of program execution
- Have to execute an **interrupt handler**

Threads

- A process may have multiple indepedent control flows called threads
- Each thread has its own stack and registers (PC, SP, registers), but share the same address space
- Shared memory model and Shared memory architecture

Architecture

L3: Processor and memory organization

L7: Cache coherence and memory consistency

L11: Interconnection networks

Parallel Computation Models

L4: Shared-memory programming models

L6: Data parallel models (GPGPU)

L9,10: Distributed-programming models

Performance and Scalability of Parallel Programs

L5: Performance of parallel systems

L8: performance instrumentation

New Trends

L12: Energy efficient computing