

L4: Query Evaluation - Sort, Select

External Merge Sort - Bocked I/O

- Read and write in blocks of **b** buffer pages (replace b with 1 for unoptimised)
- $\lfloor \frac{B-b}{b} \rfloor$ blocks for input, 1 block for output
- Can merge at most $\lfloor \frac{B-b}{b} \rfloor$ sorted runs in each merge pass
- $F = \lfloor \frac{B}{b} \rfloor - 1$ runs can be merged at each pass
- Num passes = $\log_F N_0$
- New cost: $2N(\lceil \log_F N_0 \rceil + 1)$

B+ tree sort

- B+ Tree is sorted by key
- Format 1 (clustered): Sequential Scan
- Format 2/3: Retrieve data using RID for each data entry
- Unclustered implies more I/Os

Access Path refers to the different ways to retrieve tuples from a relation. It is either a **file scan** or a **index plus matching selection condition**. The more **selective** the access paths, the fewer pages are read from the disk.

- Table scan: scan all data pages
- Index scan: scan all index pages
- Table intersection: combine results from multiple index scans (union, intersec). Find RIDs of each predicate and get the intersection

Query: Selection Covering Index

- I is a covering index of *queryQ* if I contains all attributes of Q
- No RID lookup is needed, Index-only plan
- If data is unclustered, unsorted, no index \rightarrow best way is to collect all entries and sort by RID before doing I/O

CNF Predicate

- Find RIDs of each predicate and get the intersection
- Conjuncts are in the form (R.A op c V R.a op R.b)
- CNF are conjuncts (or terms) connected by \wedge

Matching Predicates - B+ Tree

- Non-disjunctive CNF (no \vee)
- At most one non-equality comparison operator which must be on the **last attribute in the CNF**
- $(k_1 = c_1) \wedge (k_2 = c_2) \wedge \dots k_i \text{opc}_i | I = (k_1, k_2 \dots k_n)$
- The order of k matters, and there cannot be missing K_i in the middle of the CNF
- Having inequality operator before equality operator makes the query to be less selective

Matching Predicates - Hasing

- No inequality operators
- $(k_1 = c_1) \wedge \dots k_i = c_n | I = (k_1, k_2 \dots k_n)$
- Unlike B+ tree, **all predicates must match**

$I = (\text{age, weight, height})$, $p = (\text{age} \geq 20 \wedge \text{age} \geq 18 \text{weight} = 50 \wedge \text{height} = 150 \wedge \text{level} = 3)$

Primary Conjuncts : The subset of conjuncts in p that I matches

Primary Conjuncts: $\text{age} \geq 20 \wedge \text{age} \geq 18$

Covered Conjuncts : The subset of conjuncts in p that I covers (conjuncts that appear in I). Primary conjunct \subseteq covered conjunct

Covered Conjuncts: $\text{age} \geq 20 \wedge \text{age} \geq 18 \wedge \text{height} = 150$

Cost Notation

| Notation | Meaning |
|---|--|
| r | relational algebra expression |
| $ r $ | number of tuples in output of r \rightarrow data records |
| $ r $ | number of pages in output of r |
| b_r | number of data records that can fit on a page |
| b_d | number of data entries that can fit on a page |
| F | average fanout of B ⁺ -tree index (i.e., number of pointers to child nodes) |
| h | height of B ⁺ -tree index (i.e., number of levels of internal nodes) |
| $b = \lfloor \log_2(\lceil \frac{ R }{b_i} \rceil) \rfloor$ | if format-2 index on table R |
| B | number of available buffer pages |

Cost of B+-tree index evaluation of p

Let p'=primary conjuncts of p (**matching**) — p_c =covered conjuncts of p

- Navigate internal nodes to locate the first leaf page

$$Cost_{internal} = \begin{cases} \lceil \log_F(\lceil \frac{|R|}{b_d} \rceil) \rceil |Format1 \\ \lceil \log_F(\lceil \frac{|R|}{b_i} \rceil) \rceil |Otherwise \end{cases}$$

- This is traversing the height of B+ tree
- Scan leaf pages to access all qualifying data entries

$$Cost_{leaf} = \begin{cases} \lceil \frac{|\sigma_{p'}(R)|}{b_d} \rceil |Format1 \\ \lceil \frac{|\sigma_{p'}(R)|}{b_i} \rceil |Otherwise \end{cases}$$

- This is the cost of reading qualifying conjuncts
- Using p_c would be wrong since covering conjuncts may be non-matching which results in more reads from the leaves.
- Conversely, non-matching (but covered) conjuncts cannot be derived from the B+ tree and needs to be read from disk
- Retrieve qualified data records using RID lookups. 0 if I is covering OR format 1 index. $|\sigma_{p_c}(R)|$ otherwise
- includegraphics[width=5cm, height=1.3cm]optimisation.png

Cost of Hash index evaluation of p (all covered index are matched)

- Format 1: cost to retrieve **data entries** is at least $\lceil \frac{|\sigma_{p'}(R)|}{b_d} \rceil$
- Format 2: cost to retrieve **data entries** is at least $\lceil \frac{|\sigma_{p'}(R)|}{b_i} \rceil$
- Format 2: Cost to retrieve **data records** is 0 if it is a covering index (all information in data entry) OR $|\sigma_{p'}(R)|$ otherwise

L5: Query Evaluation - Projection and Join

- $\pi^*(R)$ refers to projection without removing duplicates
- $\pi(R)$ involves 1. Removing unwanted attributes 2. Removing duplicates
- Sorting is better if we have many duplicates or if hte distribution is nonuniform(overflow more likely for hashing partitions)
- Sorting allows results to be sorted
- If $B > \sqrt{|\pi_L^*(R)|}$, then both sorting and hashing has similar I/O costs ($\lceil \frac{|R|}{B} \rceil \rightarrow |R| + 2 * |\pi_L^*(R)|$)

Approach 1: project based on sorting

- Naive**: Extract attributes L from records $\rightarrow \pi_L^*(R) \rightarrow$ Sort attributes \rightarrow Remove duplicates
- Cost: Cost to scan records ($|R|$) + Cost to output to temporary result ($|\pi_L^*(R)|$) \rightarrow cost to sort records ($2|\pi_L^*(R)| \log_m(N_0) + 1$) \rightarrow Cost to scan data records ($|\pi_L^*(R)|$)

- Optimisation**: Create Sorted runs with attributes L only (Pass 0) \rightarrow Merge sorted runs and remove duplicates $\rightarrow \pi_L(R)$

Approach 2: project based on hashing

- Build a main-memory hash table to detect and remove duplicates. Insert to the hashtable if then entry is not already in it.
- Partition R into $R_1, R_2 \dots R_{B-1}$, hash on $\pi_L(t)$ for $t \in R \leftarrow (\pi_L^*(R_i))$ does not intersect $\pi_L^*(R_j), i! = j$
- Use 1 buffer for input and (B-1) for output
- Read R 1 page at a time, and hash tuples into B-1 partitions
- Flush output buffer to disk when full
- Eliminate duplicates in each partition $\pi_L^*(R_i)$
- $\pi_L(R) = \cup_{i=1}^{B-1} (\pi_L(R_i))$
- For each partition, Initialise an in-memory hash table and insert each tuple into B_j if $t \notin B_j$

Parition overflow: Hash table $\pi_L^*(R_i)$ is larger than available memory buffers.

Solution: Recursively apply hash-based partitioning to overflowed partitions.

Analysis: Effective (no overflow) when B

$$> \frac{|\pi_L^*(R)|}{B-1} * f \approx \sqrt{f * \pi_L^*(R)}$$

If no partition overflow: (partition) $|R| + \pi_L^*(R)|$ + (duplicate elimination) $|\pi_L^*(R)|$

Join

$R \bowtie_{\theta} S$, where R is the outer relation and S is the inner relation

Optimal join

- Cost: $|R| + |S|$
- load smaller relation into memory
- requires: $|S| + 2$ buffers

Tuple-based

- Cost: $|R| + |R| * |S|$
- for each tuple r in R
- for each tuple s in S
- if (r matches s) then output (r, s) 4 to result

Page-based

- Load P_R and P_S to main memory
- Cost: $|R| + |R| * |S|$
- for each page P_R in R
- for each page P_S in S
- for each tuple $r \in P_R$
- for each tuple $s \in P_S$
- if (r matches s) then output (r, s) 4 to result

Block nested-loop

- Allocate 1 page for S, 1 for output, B-2 for R $|R| \leq |S|$
- Cost: $|R| + (\lceil \frac{|R|}{B-2} \rceil * |S|)$
- $|R| \leq |S|$
- while Scanning R
- read next (B-2) pages of R to buffer
- for P_S in S
- read P_S into Buffer
- for $r \in \text{buffer} \wedge s \in P_S$
- if (r matches s) then output (r, s) 4 to result
- Without materialisation: $\lceil \frac{|R|}{B-2} \rceil * |T|$

Index Nested Loop Join

- There is an index on the join attributes of S
- Uniform distribution: r joins $\lceil \frac{|R|}{|\pi_{B_j}(S)|} \rceil$ tuples in S

- format 1 B+Tree: $|R| + |R| * J$
- Assuming unclustered:
- J = height of tree + reading leaf pages + RID Look up
- $J = \log_F(\lceil \frac{|\pi_L^*(S)|}{b_d} \rceil)$ (tree traversal) + $\lceil \frac{|\pi_{p'}(S)|}{b_i} \rceil$ (search leaf nodes) + RID lookup
- for $r \in R$
- use r to probe S's index to find matching tuples

Sort-Merge Join

- Sort R and S on join attributes and merge
- Cost: $2|\pi_L^*(R)|(\log_{B-1}(\lceil \frac{|R|}{B} \rceil) + 1) + 2|\pi_L^*(S)|(\log_{B-1}(\lceil \frac{|S|}{B} \rceil) + 1) + |\pi_L^*(R)| + |\pi_L^*(S)|$
- merging cost is $|R| + |R| * |S|$ if each tuple of R requires a full scan of S
- Optimisation: $B \geq N(R, i) + N(S, i) + 1 \geq \sqrt{|R| + |S|}$
- We can choose which relation to partition again if this is not met
- Cost: cost of getting R, S + k(write out ($|R| + |S|$)) + m(merge ($|R| + |S|$))
- If sorted on join column: $|R| + |S|$

Grace Hash Join

- Partition into $B - 1$ partitions
- If no partition overflow ($B > \sqrt{f * |S|}$):
- k(Cost to partition R, S) + Cost of probe phase
- Partition cost = cost of getting R + cost to write partitions($|R|$)
- Probe cost = $|R| + |S|$

L6: Query Optimisation

1. Search space: Queries considered

- Search Place** Queries being considered
- Linear** if at least one operand for each join is a base relation, bushy otherwise
- Left-deep** if every right join operand is a base relation
- Right-deep** if every left join operand is a base relation
- Plan enumeration - for joins between 2 tables**
- Basic DP is not always Optimal since it goes for lowest current cost and ignores the sorted property of outputs
- Enhanced DP**: left-deep only, avoids cross products, considers early selection and projections, considers order of output

3. Cost Model

- Uniformity**: uniform distribution of all values
- Independence**: Independent distribution of values in different attributes
- Inclusion**: for $R \bowtie_{\theta} S, if ||\pi_A(R)|| \leq ||\pi_B(S)||$ then $\pi_A(R) \subseteq \pi_B(S)$

Plans-no join, 1 table

- Table scan** Scan the entire table. Cost: $|R|$
- Index scan** Scan the index. Cost: $2 + |\text{leaf pages satisfying the predicate}| + |\text{entries satisfying predicate}|$ (unclustered)
- Index intersection with $I_a I_b$** Cost to find relevant entries from index and materialise(R and s) + cost to intersect partitions 1,2 (block nested, grace hash, sort merge)+ cost to RID lookup (if more attributes are needed)
- cost to partition: Scan index for matching pages + cost to write partitions from matching entries

Histogram

- Equiwidth** Each bucket has equal number of values

- Estimate: $\frac{1}{|bucket|} * ||bucket||$
- Equidepth** Each bucket has equal number of tuples
- Sub-ranges can overlap, tuples of the same value can be in 2 adjacent buckets
- $\frac{1}{|bucke_A|} * ||bucket_A|| + \frac{1}{|bucke_B|} * ||bucket_B|| + \dots$
- MCV** Separately track the top-k MCV and exclude them from the bucket

Size of query

- Join** $||R|| * ||S|| * \frac{1}{max(|\pi_b(R)|, |\pi_b(S)|)}$
- Select - OR** $(1 - (p(a! = x) * p(b! = y))) * ||R||$
- Select - AND** $p(a = x) * p(b = y) * ||R||$

L7: Transaction Management

View Equivalent

- If R_i reads A from one write W_j in S, then R_i must also read A from the same write W_j in S'
- For each data object A, Xact (if any) that performs final write on A in S must also perform final write on A in S'


Conflicting actions - WW, WR

- Dirty Read-WR** T2 read uncommitted write from T1. $W_1(X), R_2(X)$
- Unrepeatable Read-WR** T2 updates an object that T1 has previously read and T2 commits while T1 is still in progress \rightarrow T1 can get a different value from read. $R_1(X), W_2(X), C_2, R_1(X)$
- Lost Update-WW** T2 overwrites the value of an object that has been modified by T1 while T1 is still in progress. $R_1(X), R_2(X), W_1(X), W_2(X)$
- View serializable - view equivalent to some serial S** cannot be view serializable if the above anomalies occur. Conversely, no anomalies occur if view serializable.
- Blind write** $R_1(X), W_2(Y), W_1(X)$ Blind write: $W_2(Y)$
- Conflict Serializable** Conflict equivalent to serial schedule, view serializable and not blind write
- Non Conflict Serializable** find conflicting action pairs(R1(x) W2(x)), (R2(x) W1(x))
- Conflicting actions does not mean not serializable, there needs to be a cycle
- $CSS \subsetneq VSS \subsetneq MVSS$. The more restrictive, the less concurrent and less resources are needed to check serializability

Conflict Serializability Graph

- V contains a node for each committed Xact in S
- E contains (T_i, T_j) if an action precedes and conflicts with one of T_j 's actions

$R_1(A), W_2(A), Commit_2, W_1(A), Commit_1, W_3(A), Commit_3$



ACID

- Atomicity** Either all actions of a transaction are committed or none are
- Consistency** Each transaction is consistent and DB begins in a consistent state \rightarrow DB ends in a consistent state
- Isolation** Execution of one xact is isolated from other Xacts
- Durability** Once a transaction is committed, its effects persists

- Concurrency control ensures isolation
- Recovery manager ensures atomicity and durabilit
- Consistency is ensured by constraints, cascades and triggers

Schedules

- Cascading aborts** T_i read from $T_j \rightarrow T_j$ aborts $\rightarrow T_i$ aborts. This requires high book keeping efforts, so we turn to recoverable schedules instead.
- Recoverable** $\forall T \in S$ T2 must commit after T1 if T2 reads from T1 (or T2 aborts before T1). Can still have cascading aborts. Is compulsory.
- Cascadeless** Whenever T_i reads from T_j in S, Commit must precede this action. Desirable, not compulsory.
- Theorem 4: Cascadeless \rightarrow Recoverable (not iff)
- Recovery with before image** Store the before image before write, restore this before image if write aborts. This can lead to **lost update** anomaly if the before image overwrites another Xact's write.
- Strict** to use before-images, $\forall W_i(O) \in S$, O is not read or written by another Xact until Ti either aborts or commits. This ensures no lost update anomaly during recovery.
- Strict schedules allows recovery using before images to be more efficient but restricts concurrency
- Theorem 5: Strict \rightarrow Cascadless (not iff)

L8: Concurrency Control

Lock based concurrency control

- If lock request is not granted, then T becomes blocked and gets added to O's request queue

2PL

- To read an object O, a Xact must hold a S-lock or X-lock on O
- To write to an object O, a Xact must hold a X-lock on O
- Once a Xact releases a lock, the Xact can't request any more locks
- Theorem 1: 2PL is conflict serializable

Strict 2PL

- A Xact must hold on to locks until Xact commits or aborts
- Theorem 2: Strict 2PL is strict and conflict serializable
- Strict 2PL prevents cascading rollback and deadlock and ensures recoverability

Anomalies not dealt by strict 2PL

- Phantom Read:** T1 reads a set of objects, T2 inserts a new object in that set, T1 reads the set again and gets a different set of objects. 2PLS cannot prevent this as locks are held at the object level.
- Solved by **predicate locking**, which is done in practice by **index locking**

| Isolation Level | Dirty Read | Unrepeatable Read | Phantom Read |
|------------------|--------------|-------------------|--------------|
| READ UNCOMMITTED | possible | possible | possible |
| READ COMMITTED | not possible | possible | possible |
| REPEATABLE READ | not possible | not possible | possible |
| SERIALIZABLE | not possible | not possible | not possible |

Deadlock Detection

- Waits-for graph (WFG) \rightarrow Deadlock is detected if WFG has a cycle. $(V_i, V_j \rightarrow T_i waits - for T_j)$
- Breaks a deadlock by aborting a Xact in cycle
- Alt: timeout mechanism

Deadlock Prevention

- Each Xact is assigned a timestamp when it starts

- Assume older (smaller time stamp) Xacts have higher priority than younger Xacts
- Tie between blocked/restarted xact brokered by priority, original timestamp is maintained to prevent starvation
- Wait-die** Higher priority waits for lower priority, lower priority dies if higher priority holds lock (lower never waits)
- Wound-wait** Higher priority kills lower priority, lower priority waits for higher priority to release lock (higher never waits)

Lock Conversion

- Allows for greater concurrency
- Conversion is only allowed if the Xact has not released any lock
- Upgrade(A)** blocked if another Xact holds shared or exclusive lock on A
- Downgrade(A)** allowed if Xact has not modified A and Xact has not released any lock

Improve System Throughput

- Reduce Lock Granularity, Reduce time of lock being held, reduce hotspots in DB by changeinge schema design

Multi Version Serializable Schedle (MVSS)

Benefits

- $W_i(O)$ Creates new version
- Read-only Xact are not blocked by update xact (vice versa). Mono version (e.g. 2PL) blocks
- Read-only xacts are never aborted (no deadlocks due to Multi-version) or blocked
- multiversion view equivalent** if S and S' have the same set of read-from relationships
- i.e. $R_i(x_j)$ occurs in S iff $R_i(x_j)$ occurs in S'
- Monoversion Schedule** each read action returns the most recently created object version. Not necessarily serializable
- MVSS** if there exists a serial Monoversion schedule that is multiversion view equivalent to S
- Note that a MVSS is not necessarily conflict serializable schedule if it is not a valid monoversion schedule
- E.g. $W1(x1), R2(x0), R2(y0), W1(y1), C1, C2$ is MVSS with (T2, T1) but contains conflicting actions $W1(x1)$ and $R2(X0)$

Snapshot Isolation (SI) [NOT always serializable]

- Similar performance as Read Committed, does not suffer lost update, unrepeatable read.
- Each Xact has a snapshot of the database at the start of the Xact and sees only versions from that snapshot and **its own writes**
- Concurrent Update Property** If multiple concurrent exacts on same object, only one can commit
- FUW** T needs to acquire X-lock on O (if not - wait), and if O has been updated by a concurrent T' then T aborts
- FCW** (no locks) before committing T checks if O has been updated, abort if it has been updated
- Write-skew anomaly**, not MVSS: $R_1(X_0), R_2(X_0), R_1(Y_1), R_2(Y_2), W_1(X_1), C_1, W_2(Y_2), C_2$
- Read-only anomaly**,not MVSS: $R_1(b), R_2(a), W_1(b), C_1, R_2(b), W_2(a), R_3(a), R_3(b), C_3, C_2$
- Serializable** MVSS to some monoversion

Transaction Dependencies - Making SI serializable

- WW** from T1 to T2: T1 commits some version of X and T2 writes the immediate successor

- WR** from T1 to T2: T1 commits some version of X which is read by T2
- RW** from T1 to T2: T1 reads some version of X and T2 commits the immediate successor
- DSG** V = xacts, E = Dependencies, use \rightarrow for concurrent transactions and \rightarrow for non-concurrent
- Non-MVSS SI** At least one cycle in DSG(s) with T_i, T_j, T_k s.t T_i, T_k are possibly the same xact, T_i, T_j are concurrent with an edge T_i rw T_j and T_j rw T_k

Locking Granularity

(Most coarse) DB, relation, page, tuple (Finest, more concurrency, less scalable)

Locks are acquired in a top down manner.

| Lock Requested | Lock Held | | | | |
|----------------|-----------|----|----|---|---|
| | - | IS | IX | S | X |
| IS | ✓ | ✓ | ✓ | ✓ | × |
| IX | ✓ | ✓ | ✓ | ✓ | × |
| S | ✓ | ✓ | × | ✓ | × |
| X | ✓ | × | × | × | × |

L9-Crash Recovery

Policies

- Steal:** Allows dirty pages to write to disk before commit
- No Steal reduces number of free buffer pages
- Force:** All dirty pages must write to disk when commit
- Force incurs more random disk IO
- Steal: undo, Force: no redo. Pgsq: steal and no-force

Restart: analysis, redo, undo

- Analysis: identifies dirtied pool pages and active Xacts at time of crush
- Redo: redo actions to restore db to pre-crush
- Undo: undo actions of Xacts that did not commit

Analysis: Xact table

- When the first log record is created, create a new entry T with status U
- Update lastLSN for T to be r's LSN
- Remove T if end log is seen

Analysis: Dirty Page Table

- New dirtied page will be added to the DPT with recLSN=r.LSN
- Remove entry when it is flushed to disk

| | | LOG | | | | | | | | |
|----|---------|----------------|--------|--------|--------|--------|--------------|-------------|--|--|
| | prevLSN | XactID | type | pageID | length | offset | before image | after image | | |
| 10 | - | T ₁ | update | P500 | 3 | 21 | ABC | DEF | | |
| 20 | - | T ₂ | update | P600 | 3 | 41 | HJ | KLM | | |
| 30 | 20 | T ₂ | update | P500 | 3 | 20 | GDE | QRS | | |
| 40 | 10 | T ₁ | update | P505 | 3 | 21 | TUV | WXY | | |

| DIRTY PAGE TABLE | | | XACT TABLE | | |
|------------------|--------|--|----------------|---------|--------|
| pageID | recLSN | | XactID | lastLSN | status |
| P500 | 10 | | T ₁ | 40 | U |
| P600 | 20 | | T ₂ | 30 | U |
| P505 | 40 | | | | |

Redo Phase (DPT)

- Redo LSN = min(recLSNs), then fetch page LSN,
- if r.LSN > pageLSN and Page is in DPT, redo
- Update pageLSN to r.LSN

Undo Phase (TT)

- Start from largest LSN from L
- if update, create CLR with undoNextLSN=r's prevLSN, update-L-TT(r.prevLSN)
- if CLR, update-L-TT(r.undoNextLSN)
- update-L-TT(Isn): add Isn to L if Isn not null, else add end log record for T and remove it from TT

Checkpointing

- Normal(no ECPLR): CPLR's TT, empty DPT
- Fuzzy: BeginCPLR's TT and BCPLR's DPT