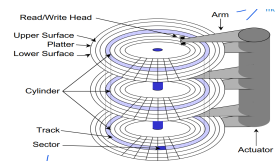


## L1 - Data Storage

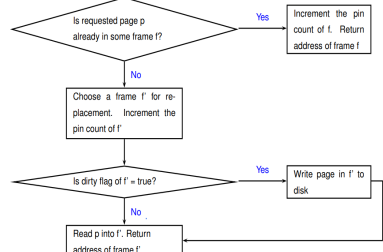
### Magnetic Disks



- Disk Access Time** Seek time + Rotational Latency + Transfer time
- Response time** Queueing delay + Disk access time
- Rotational Delay**  $\frac{1}{2} \frac{60s}{RPM}$
- Transfer Time** sectors on the same track \*  $\frac{TimePerRevolution}{SectorsPerTrack}$

### Buffer Manager

- Buffer pool** Main memory allocated for DBMS
- pin count** is incremented upon pinning
- dirty bit** is updated when the page is unpinned (if modified)
- Replacement is only possible if pin count == 0



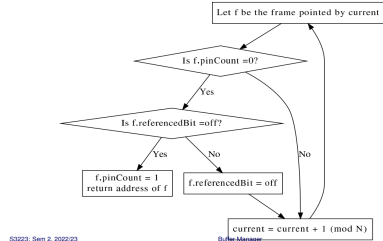
### Replacement Policies

#### LRU Policy

- Maintains a queue of pointers to frames with pin count = 0

#### Clock Replacement Policy

N = number of frames in buffer pool

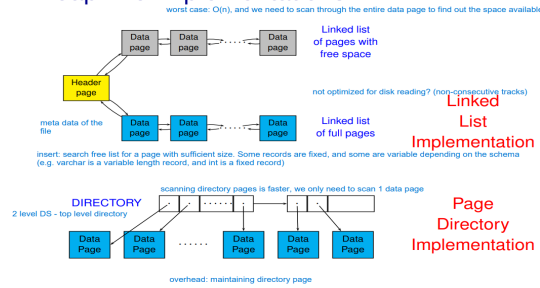


CS3223, Sem 2, 2022/23

- Simplifies LRU with a second chance round robin system
- Each frame has a **reference bit** that is turned on when pin count reaches 0
- Replaces a page when referenced bit is off and pin count is 0

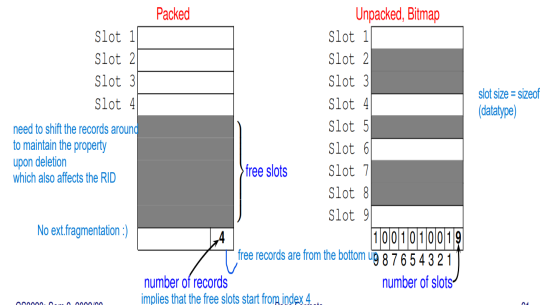
## File Organisation

### Heap File Implementations



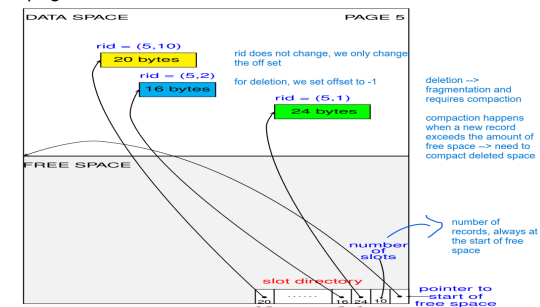
### Page Formats: Fixed Length Records

- Packed Organisation** Store records in contiguous slots
- Unpacked Organisation** Uses a bit array to maintain free slots



### Page Formats: Slotted Page (variable length record)

- Store records in slots of (record offset, record length)
- Record Offset:** Offset of the record from the start of the page



### Record Formats

#### Fixed-Length Records

- Fields are stored consecutively

F1 | F2 | F3 | F4

#### Variable-Length Records

- Delimit fields with special symbols

F1 | \$ | F2 | \$ | F3 | \$ | F4

- Use an array of field offsets

o1 | o2 | o3 | o4 | F1 | F2 | F3 | F4

- Each o<sub>i</sub> is an offset to beginning of field F<sub>i</sub>

## L2 - Indexing

- A search key is a sequence of k attributes. If  $k \geq 1$ , composite key
- A search key is a unique index if it is a candidate key
- An index is stored as a file
- Format of data entries**
- Format 1:  $k^*$  is an actual data record with search value k

- Format 2:  $k^*$  is the form (k, rid)
- Format 3:  $k^*$  is the form (k, rid-list\*)
- Note: Different formats affects the number of data entries stored in a page

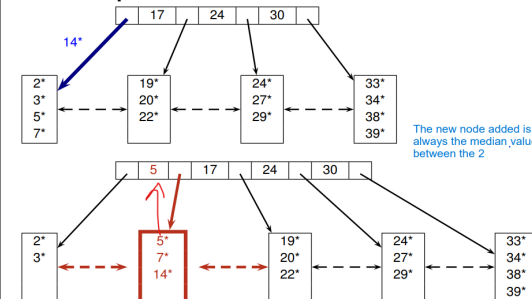
### Clustered Vs Unclustered

- Clustered:** Order of data entries is the same as the order of data records. Can only be built on ordered field (e.g. primary key)
- Unclustered:** Order of data entries does not correspond to the order of data records
- The implication is that we can read an entire clustered page with 1 I/O
- B+ Tree: Format 1 is clustered, Format 2 and 3 can be clustered if data records are sorted on the search key
- Hash: Only format 1 is clustered since hashing does not store data entries in search key order

### Tree Based Index - B+ Tree

- Leaf nodes are doubly linked and store Data Entries
- Internal nodes store index entries (p<sub>0</sub>, k, p<sub>1</sub> ... p<sub>k</sub>, k, p<sub>k+1</sub>)
- Internal nodes contains m entries,  $m \in [d, 2d]$  → space utilisation  $\geq 50\%$
- Root contains m entries,  $m \in [1, 2d]$



### B+ Tree - Split Overflow Nodes





- Distribute d+1 entries to the new leaf node
- Create new entry index using smallest key in the new node (middle key)
- Insert new entry into parent node of overflowed node

### B+ Tree - Overflow Propagation


5 is pushed up



17 is pushed up





- Excess middle node is pushed updated to parent node
- B+ Tree - Redistribution of data entries**
- Two nodes are siblings if they have the same parent node



### B+ Tree - Underflow

- Underflow occurs when a node has less than d entries
- Underflow is resolved by redistributing entries between siblings
- An underflow node is merged if each of its adjacent siblings have exactly d entries



### B+ Tree - Bulk Loading

- Initializing a B+ tree by insertion is expensive (need to traverse tree n times)
- 1. Sort all data entries by search key
- 2. Initialize B+ tree with an empty root page
- 3. Load data entries into leaf pages
- 4. In asc order, insert the index entry of each leaf page into the rightmost parent node

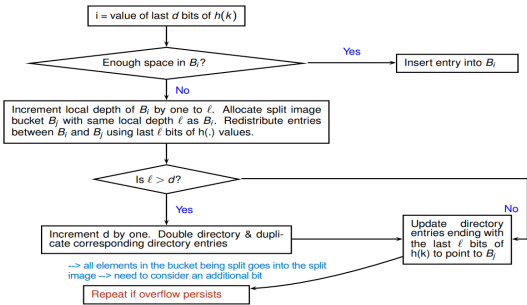
### Hash Based Index

- Does not support range search, only equality queries
- Static Hashing**

- N buckets, each bucket has 1 primary page and  $\geq 0$  overflow pages
- To maintain performance, we need to routinely construct bigger hash tables and redistribute data entries

Dynamic Hashing - Extensible Hashing

- No overflow pages! A bucket can be thought of as a page
- At most 2 Disk I/Os for equality search (at most 1 if directory and bucket fits in memory)
- Instead of maintaining data entries, we maintain pointers to data entries in buckets
- Instead of maintaining buckets, maintain a directory of pointers to buckets
- The directory has  $2^d$  buckets, where d is the global depth  $\rightarrow$  large overhead if hashing is uniform
- Each director entry differs by a unique d-bit address
- Two directories are corresponding iff their addresses differ only in the dth bit
- All entries with the same local depth (l) have the same last l bits in h(k)



Extensible Hashing - Split, Double


- Split and doubling is checked every time a bucket is full
- Doubling only happens if local depth = global depth
- The split image has the same depth as the split bucket
- Other than the split image of the split bucket, split image of other buckets points to the same corresponding bucket
- Each bucket is pointed by  $2^{(d-l)}$  directories

Extensible Hashing - Deletion

- $B_i$  is deallocated

- l decrement by 1
- Directory Entries that point to  $B_i$  points to its corresponding bucket

Dynamic Hashing - Linear Hashing



- **GetBucketNum(k)** returns bucket # where entry with search key  $k$  is located
- $$GetBucketNum(k) = \begin{cases} h_{level}(k) & \text{if } h_{level}(k) \geq next, \\ h_{level+1}(k) & \text{otherwise.} \end{cases}$$
- **SplitBucket()** splits bucket  $B_{next}$
- 1. Redistribute the entries in  $B_{next}$  into  $B_{next+N_{used}}$  using  $h_{level+1}()$
  2.  $next = next + 1$
  3. if  $(next == N_{used})$  then { level = level + 1; next = 0 }

- One I/O for equality search (more per number of overflow pages in bucket)
- Performs worse than extensible hashing if distribution is skewed

- Does not require a directory
  - Higher average space utilisation, but longer overflow chains
  - Has a family of hash functions, with each having a range twice of its predecessor
  - $N_0$ : initial number of buckets
  - $N_i = 2^i N_0$ : number of buckets at start of round i
  - $next$ : the next bucket to be split, this is incremented every time split happens
  - $h_i = h(k) \bmod N_i$ : hash function for round i, if the bucket  $\leq next$  (already split)
  - $h_{i+1} = h(k) \bmod N_{i+1}$ : hash function for round i+1, if the bucket  $> next$
  - Split Criteria: By default, split when a bucket overflows
- Linear Hashing - Deletion**
- Essentially the inverse of insertion
  - If the last bucket is empty  $\rightarrow$  delete it and decrement  $next$  by 1
  - If  $next$  is 0, set it to  $M/2 - 1$ , and we can decrement level by 1 (half of buckets have been deleted if  $next$  is 0)
  - Merging with corresponding bucket is optional