

L4: Query Evaluation - Sort, Select

Sorting - External Merge Sort

Projection, join, bulk loading etc all require sorting

► Uses B number of buffer pages

► **Pass 0:** Creation of sorted runs

- Read in and sort B pages at a time
- Number of sorted runs created = $\lceil N/B \rceil$
- Size of each sorted run = B pages (except possibly for last run)

► **Pass i, $i \geq 1$:** Merging of sorted runs

- Use $B-1$ buffer pages for input & one buffer page for output
- Performs (B-1)-way merge

► **Analysis:**

- N_0 = number of sorted runs created in pass 0 = $\lceil N/B \rceil$
- Total number of passes = $\lceil \log_{B-1}(N_0) \rceil + 1$
- Total number of I/O = $2N(\lceil \log_{B-1}(N_0) \rceil + 1)$

★ Each pass reads N pages & writes N pages

External Merge Sort - Buffered I/O

- Read and write in blocks of **b** buffer pages (replace b with 1 for unoptimised)
- $\lfloor \frac{B-b}{b} \rfloor$ blocks for input, 1 block for output
- Can merge at most $\lfloor \frac{B-b}{b} \rfloor$ sorted runs in each merge pass
- $F = \lfloor \frac{B}{b} \rfloor - 1$ runs can be merged at each pass
- Num passes = $\log_F N_0$

B+ tree sort

- B+ Tree is sorted by key
- Format 1 (clustered): Sequential Scan
- Format 2/3: Retrieve data using RID for each data entry
- Unclustered implies more I/Os

Access Path refers to the different ways to retrieve tuples from a relation. It is either a **file scan** or a **index plus matching selection condition**. The more **selective** the access paths, the fewer pages are read from the disk.

- Table scan: scan all data pages
- Index scan: scan all index pages
- Table intersection: combine results from multiple index scans (union, intersec). Find RIDs of each predicate and get the intersection

Query: Selection Covering Index

- I is a covering index of *query*_Q if I contains all attributes of Q
- No RID lookup is needed, Index-only plan
- If data is unclustered, unsorted, no index - \hat{c} best way is to collect all entries and sort by RID before doing I/O

CNF Predicate

- Find RIDs of each predicate and get the intersection
- Conjuncts are in the form (R.A op c V R.a op R.b)
- CNF are conjuncts (or terms) connected by \wedge

Matching Predicates - B+ Tree

- Non-disjunctive CNF (no \vee)
- At most one non-equality comparison operator which must be on the **last attribute in the CNF**
- $(k_1 = c_1) \wedge (k_2 = c_2) \wedge \dots k_i \text{opc}_i | I = (k_1, k_2 \dots k_n)$
- The order of k matters, and there cannot be missing K_i in the middle of the CNF

- Having inequality operator before equality operator makes the query to be less selective

Matching Predicates - Hasing

- No inequality operators

$$(k_1 = c_1) \wedge \dots k_i = c_n | I = (k_1, k_2 \dots k_n)$$

- Unlike B+ tree, **all predicates must match**

$I = (\text{age, weight, height})$, $p = (\text{age} \geq 20 \wedge \text{age} \geq 18 \text{weight} = 50 \wedge \text{height} = 150 \wedge \text{level} = 3)$

Primary Conjuncts : The subset of conjuncts in p that I matches

Primary Conjuncts: $\text{age} \geq 20 \wedge \text{age} \geq 18$

Covered Conjuncts : The subset of conjuncts in p that I covers (conjuncts that appear in I). Primary conjunct \subseteq covered conjunct

Covered Conjuncts: $\text{age} \geq 20 \wedge \text{age} \geq 18 \wedge \text{height} = 150$

Cost Notation

Notation	Meaning
r	relational algebra expression
$ r $	number of tuples in output of r <small>data records</small>
$ r $	number of pages in output of r
b_d	number of data records that can fit on a page
b_i	number of data entries that can fit on a page
F	average fanout of B ⁺ -tree index (i.e., number of pointers to child nodes)
h	height of B ⁺ -tree index (i.e., number of levels of internal nodes)
	$h = \lceil \log_F(\lceil \frac{ r }{b_i} \rceil) \rceil$ if format-2 index on table R
B	number of available buffer pages

Cost of B+-tree index evaluation of p

Let p' = primary conjuncts of p — p_c = covered conjuncts of p

1. Navigate internal nodes to locate the first leaf page

$$Cost_{internal} = \begin{cases} \lceil \log_F(\lceil \frac{|r|}{b_d} \rceil) \rceil |Format 1| \\ \lceil \log_F(\lceil \frac{|r|}{b_i} \rceil) \rceil |Otherwise| \end{cases}$$

- 1.1 This is traversing the height of B+ tree
2. Scan leaf pages to access all qualifying data entries

$$Cost_{leaf} = \begin{cases} \lceil \frac{|\sigma_{p'}(R)|}{b_d} \rceil |Format 1| \\ \lceil \frac{|\sigma_{p'}(R)|}{b_i} \rceil |Otherwise| \end{cases}$$

- 2.1 This is the cost of reading qualifying conjuncts
- 2.2 Using p_c would be wrong since covering conjuncts may be non-matching which results in more reads from the leaves
- 3 Retrieve qualified data records using RID lookups. 0 if I is covering OR format 1 index. $|\sigma_{p_c}(R)|$ otherwise

Cost of RID lookups could be reduced by first sorting the RIDs

$$\lceil \frac{|\sigma_{p_c}(R)|}{b_d} \rceil \leq Cost_{rid} \leq \min\{|\sigma_{p_c}(R)|, |R|\} \quad \text{assuming clustered/congruous} \quad \text{sequenced I/O}$$

ceiling because we have to read the additional page for the remainder RIDs

Cost of Hash index evaluation of p

- Format 1: cost to retrieve **data entries** is at least $\lceil \frac{|\sigma_{p'}(R)|}{b_d} \rceil$
- Format 2: cost to retrieve **data entries** is at least $\lceil \frac{|\sigma_{p'}(R)|}{b_i} \rceil$
- Format 2: cost to retrieve **data records** is 0 if it is a covering index (all information in data entry) OR $|\sigma_{p'}(R)|$ otherwise

L5: Query Evaluation - Projection and Join

- $\pi^*(R)$ refers to projection without removing duplicates

- $\pi(R)$ involves 1. Removing unwanted attributes 2. Removing duplicates
- Sorting is better if we have many duplicates or if the distribution is nonuniform (overflow more likely for hashing partitions)
- Sorting allows results to be sorted
- If $B > \sqrt{|\pi_L^*(R)|}$, then both sorting and hashing has similar I/O costs ($\lceil \frac{|R|}{B} \rceil \rightarrow |R| + 2 * |\pi_L^*(R)|$)

Approach 1: project based on sorting

- **Naive:** Extract attributes L from records $\rightarrow \pi_L^*(R) \rightarrow$ Sort attributes \rightarrow Remove duplicates
- Cost: Cost to scan records ($|R|$) + Cost to output to temporary result ($|\pi_L^*(R)|$) \rightarrow cost to sort records ($2|\pi_L^*(R)| \log_m(N_0) + 1$) \rightarrow Cost to scan data records ($|\pi_L^*(R)|$)
- **Optimisation:** Create Sorted runs with attributes L only (Pass 0) \rightarrow Merge sorted runs and remove duplicates $\rightarrow \pi_L(R)$

Approach 2: project based on hashing

- Build a main-memory hash table to detect and remove duplicates. Insert to the hashtable if then entry is not already in it.
- 1. Partition R into $R_1, R_2 \dots R_{B-1}$, hash on $\pi_L(t)$ for $t \in R \leftarrow (\pi_L^*(R_i) \text{ does not intersect } \pi_L^*(R_j), i! = j)$
- 1.1 Use 1 buffer for input and (B-1) for output
- 1.2 Read R 1 page at a time, and hash tuples into B-1 partitions
- 1.3 Flush output buffer to disk when full
- 2. Eliminate duplicates in each partition $\pi_L^*(R_i)$
- $\pi_L(R) = \cup_i^{B-1} (\pi_L^*(R_i))$
- 2.1 For each partition, Initialise an in-memory hash table and insert each tuple into B_j if $t \notin B_j$

Partition overflow: Hash table $\pi_L^*(R_i)$ is larger than available memory buffers.

Solution: Recursively apply hash-based partitioning to overflowed partitions.

Analysis: Effective (no overflow) when B

$$> \frac{|\pi_L^*(R)|}{B-1} * f \approx \sqrt{f * \pi_L^*(R)}$$

If no partition overflow: (partition) $|R| + \pi_L^*(R)$ + (duplicate elimination) $|\pi_L^*(R)|$

Index based projection: Do index scan if the wanted attributes \subseteq search key

Join $R \bowtie_{\theta} S$, where R is the outer relation and S is the inner relation

- **Tuple-based**
 - Cost: $|R| + |R| * |S|$
 - for each tuple r in R
 - for each tuple s in S
 - if (r matches s) then output (r, s) to result
- **Page-based**
 - Load P_R and P_S to main memory
 - Cost: $|R| + |R| * |S|$
 - for each page P_R in R
 - for each page P_S in S
 - for each tuple $r \in P_R$
 - for each tuple $s \in P_S$
 - if (r matches s) then output (r, s) to result
- **Block nested-loop**
 - Allocate 1 page for S, 1 for output, B-2 for R
 - $|R| \leq |S|$
 - Cost: $|R| + (\lceil \frac{|R|}{B-2} \rceil * |S|)$
 - $|R| \leq |S|$
 - while Scanning R
 - read next (B-2) pages of R to buffer
 - for P_S in S
 - read P_S into Buffer
 - for $r \in \text{buffer} \wedge s \in P_S$
 - if (r matches s) then output (r, s) to result
 - Without materialisation: $\lceil \frac{|R|}{B-2} \rceil * |T|$
- **Index Nested Loop Join**
 - There is an index on the join attributes of S
 - Uniform distribution: r joins $\lceil \frac{|S|}{|\pi_{B_j}(S)|} \rceil$ tuples in S
 - format 1
 - B+Tree: $|R| + |R| * J$
- **J = $\log_F(\lceil \frac{|S|}{b_d} \rceil)$ (tree traversal) + $\lceil \frac{|S|}{b_d |\pi_{B_j}(S)|} \rceil$ (search leaf nodes)**
- for $r \in R$
- use r to probe S's index to find matching tuples
- **Sort-Merge Join**
 - Sort R and S on join attributes and merge
 - Cost: $2|\pi_L^*(R)|(\log_{B-1}(\frac{|R|}{B}) + 1) + 2|\pi_L^*(S)|(\log_{B-1}(\frac{|S|}{B}) + 1) + |\pi_L^*(R)| + |\pi_L^*(S)|$
 - merging cost is $|R| + |R| * |S|$ if each tuple of R requires a full scan of S
 - Optimisation: $B \geq \frac{N(R, i) + N(S, i) + 1}{\sqrt{|R| + |S|}}$
 - We can choose which relation to partition again if this is not met
 - Cost: cost of getting R, S + k(write out $(|R| + |S|)$) + m(merge $(|R| + |S|)$)
 - If sorted on join column: $|R| + |S|$
- **Grace Hash Join**
 - Partition into $B - 1$ partitions
 - If no partition overflow ($B > \sqrt{f * |S|}$):
 - k(Cost to partition R, S) + Cost of probe phase
 - Partition cost = cost of getting R + cost to write partitions($|R|$)
 - Probe cost = $|R| + |S|$

L6: Query Optimisation

1. Search space: Queries considered

- **Search Place** Queries being considered
- **Linear** if at least one operand for each join is a base relation, bushy otherwise
- **Left-deep** if every right join operand is a base relation
- **Right-deep** if every left join operand is a base relation

2. Plan enumeration - for joins between 2 tables

Dynamic programming formulation SPJ = SELECT, PROJECT, JOIN

Input: A SPJ query q on relations R_1, R_2, \dots, R_n

Output: An optimal query plan for q

NP Hard (Normally there is a limit to how many tables are joined, worst case 12, any larger we won't use DP)

```

01. for i = 1 to n do // consider 1 relation plan -> best plan for accessing that table
02.   optPlan({Ri}) = best access plan for Ri
03. for i = 2 to n do // consider all subset of relations of size 1
04.   for each S ⊆ {R1, ..., Rn}, |S| = i do // subsets
05.     bestPlan = dummy plan with cost(bestPlan) = ∞
06.     for each Sj, Sk, |Sj| ∈ [1, i], S = Sj ∪ Sk do // ways to form this subset
07.       p = best way to join optPlan(Sj) and optPlan(Sk)
08.       if (cost(p) ≤ cost(bestPlan)) then // use the optimal plan identified in previous iterations
09.         bestPlan = p
10.   optPlan(S) = bestPlan // best plan for subset is updated
11. return optPlan({R1, ..., Rn})
  
```

- Decide on the plan for the 2 operands
- Decide on the plan to join: Block nested loop, sort merge join, grace hash join

3. Cost Model

- Uniformity:** uniform distribution of all values
- Independence:** Independent distribution of values in different attributes
- Inclusion:** for $R \bowtie S, i_f || \pi_A(R) || \leq || \pi_B(S) ||$ then $\pi_A(R) \subseteq \pi_B(S)$

Plans-no join, 1 table

- Table scan** Scan the entire table. Cost: $|R|$
- Index scan** Scan the index. Cost: $2 + |\text{leaf pages satisfying the predicate}| + |\text{entries satisfying predicate}|$ (unclustered)
- Index intersection with I_a, I_b** Cost to partition predicate1(R) + Cost to partition predicate2(R) + cost to intersect partitions 1,2 + cost to RID lookup
- cost to partition: Scan index for matching pages + cost to write partitions from matching entries

Histogram

- Equiwidth** Each bucket has equal number of values
- Estimate: $\frac{1}{|\text{bucket}|} * ||\text{bucket}||$
- Equidepth** Each bucket has equal number of tuples
- Sub-ranges can overlap, tuples of the same value can be in 2 adjacent buckets
- $\frac{1}{|\text{bucket}_A|} * ||\text{bucket}_A|| + \frac{1}{|\text{bucket}_B|} * ||\text{bucket}_B|| + \dots$
- MCV** Separately track the top-k MCV and exclude them from the bucket

Size of query

- Join** $||R|| * ||S|| * \frac{1}{\max(||\pi_b(R)||, ||\pi_b(S)||)}$
- Select - OR** $(1 - (p(a = x) * p(b = y))) * ||R||$
- Select - AND** $p(a = x) * p(b = y) * ||R||$

L7: Transaction Management

View Equivalent

- If T_i reads A from T_j in S, then T_i must also read A from T_j in S'
- For each data object A, Xact (if any) that performs final write on A in S must also perform final write on A in S'

Conflicting actions - WW, WR


- Dirty Read** T2 read uncommitted write from T1
- Unrepeatable Read** T2 updates an object that T1 has previously read and T2 commits while T1 is still in progress \rightarrow T1 can get a different value from read
- Lost Update** T2 overwrites the value of an object that has been modified by T1 while T1 is still in progress
- View serializable prevents These

- Blind write** $R_1(X), W_2(Y), W_1(X)$ Blind write: $W_2(Y)$
- Conflict Serializable** Conflict equivalent to serial schedule, view serializable and not blind write
- Non Conflict Serializable** find conflicting action pairs(R1(x) W2(x)), (R2(x) W1(x))
- Conflicting actions does not mean not serializable, there needs to be a cycle
- $CSS \subsetneq VSS \subsetneq MVSS$

Conflict Serializability Graph

- V contains a node for each committed Xact in S
- E contains (T_i, T_j) if an action precedes and conflicts with one of T_j 's actions

$R_1(A), W_2(A), Commit_2, W_1(A), Commit_1, W_3(A), Commit_3$



Schedules

- Cascading aborts** T_i read from $T_j \rightarrow T_j$ aborts $\rightarrow T_i$ aborts
- Recoverable** $\forall T \in S$ T2 must commit after T1 if T2 reads from T1 (or T2 aborts before T1)
- Cascadeless** Whenever T_i reads from T_j in S, Commit must precede this action
- Theorem 4: Cascadeless \rightarrow Recoverable (not iff)
- Strict** to use before-images, $\forall W_i(O) \in S, O$ is not read or written by another Xact until Ti either aborts or commits
- Theorem 5: Strict \rightarrow Cascadless (not iff)

2PL

- To read an object O, a Xact must hold a S-lock or X-lock on O
- To write to an object O, a Xact must hold a X-lock on O
- Once a Xact releases a lock, the Xact can't request any more locks
- Theorem 1: 2PL is conflict serializable

Strict 2PL

- A Xact must hold on to locks until Xact commits or aborts
- Theorem 2: Strict 2PL is strict and conflict serializable
- Strict 2PL prevents cascading rollback and deadlock and ensures recoverability

Detect deadlocks

- Waits-for graph (WFG) \rightarrow Deadlock is detected if WFG has a cycle. $(V_i, V_j \rightarrow T_i \text{ waits for } T_j)$
- Breaks a deadlock by aborting a Xact in cycle

Deadlock Prevention

- Each Xact is assigned a timestamp when it starts

- Assume older (smaller time stamp) Xacts have higher priority than younger Xacts
- Tie between blocked/restarted xact brokered by priority, priority is maintained after abort
 - Suppose T_i requests for a lock that conflicts with a lock held by T_j
 - Two possible deadlock prevention policies:
 - Wait-die policy:** lower-priority Xacts never wait for higher-priority Xacts
 - Wound-wait policy:** higher-priority Xacts never wait for lower-priority Xacts

Prevention Policy	T_i has higher priority	T_i has lower priority
Wait-die	T_i waits for T_j	T_i aborts
Wound-wait	T_j aborts	T_i waits for T_j

L8: MVCC

Multi Version Serializable Schedle (MVSS)

- multiversion view equivalent** if S and S' have the same set of read-from relationships
- i.e. $R_i(x_j)$ occurs in S iff $R_i(x_j)$ occurs in S'
- Monoversion Schedule** each read action returns the most recently created object version
- MVSS** if there exists a serial Monoversion schedule that is multiversion view equivalent to S
- Note that a MVSS is not necessarily conflict serializable schedule if it is not a valid monoversion schedule
- E.g. $W1(x1), R2(x0), R2(y0), W1(y1), C1, C2$ is MVSS with (T2, T1) but contains conflicting actions $W1(x1)$ and $R2(X0)$

Snapshot Isolation (SI)

- Each Xact has a snapshot of the database at the start of the Xact and sees only versions from that snapshot and **its own writes**
- FUW** T needs to acquire X-lock on O (if not - wait), and if O has been updated by a concurrent T' then T aborts
- FCW** (no locks) before committing T checks if O has been updated, abort if it has been updated
- Write-skew anomaly**, not MVSS: $R_1(X_0), R_2(X_0), R_1(Y_1), R_2(Y_2), W_1(X_1), Commit1, W_2(Y_2), Commit2$
- Read-only anomaly**,not MVSS: $R_1(b), R_2(a), W_1(b), C_1, R_2(b), W_2(a), R_3(a), R_3(b), C_2, C_1$
- Transaction Dependencies**
 - WW** from T1 to T2: T1 commits some version of X and T2 writes the immediate successor
 - WR** from T1 to T2: T1 commits some version of X which is read by T2
 - RW** from T1 to T2: T1 reads some version of X and T2 commits the immediate successor

- DSG** V = xacts, E = Dependencies, use \rightarrow for concurrent transactions and \rightarrow for non-concurrent

Lock Requested	Lock Held				
	-	IS	IX	S	X
IS		✓	✓	✓	×
IX		✓	✓	✓	×
S		✓	×	×	×
X		✓	×	×	×

L10-Recovery

Policies

- Steal:** Allows dirty pages to be written to disk before commit
- Force:** Requires all dirty pages to be written to disk when commit
- No-steal: no undo, Force: no redo. Pgsql uses steal and no-force

Restart: analysis, redo, undo

- Analysis: identifies dirtied pool pages and active Xacts at time of crush
- Redo: redo actions to restore db to pre-crush
- Undo: undo actions of Xacts that did not commit

Analysis: Xact table

- When the first log record is created, create a new entry T with status U
- Update lastLSN for T to be r's LSN
- Remove T if end log is seen

Analysis: Dirty Page Table

- New dirtied page will be added to the DPT with recLSN=r.LSN
- Remove entry when it is flushed to disk

		LOG					
	prevLSN	XactID	type	pageID	length	offset	
10	-	T ₁	update	P500	3	21	ABC
20	-	T ₂	update	P600	3	41	HJ
30	20	T ₂	update	P500	3	20	GDE
40	10	T ₁	update	P505	3	21	TUV

DIRTY PAGE TABLE				XACT TABLE		
pageID	recLSN			XactID	lastLSN	status
P500	10			T ₁	40	U
P600	30			T ₂	30	U
P505	40					

Redo Phase (DPT)

- Redo LSN = min(recLSNs), then fetch page LSN,
- if r.LSN > pageLSN and Page is in DPT, redo
- Update pageLSN to r.LSN

Undo Phase (TT)

- Start from largest LSN from L
- if update, create CLR with undoNextLSN=r's prevLSN, update-L-TT(r.prevLSN)
- if CLR, update-L-TT(r.undoNextLSN)
- update-L-TT(Isn): add Isn to L if Isn not null, else add end log record for T and remove it from TT

Checkpointing

- Normal(no ECPLR): CPLR's TT, empty DPT
- Fuzzy: BeginCPLR's TT and BCPLR's DPT