


L1 - Data Storage

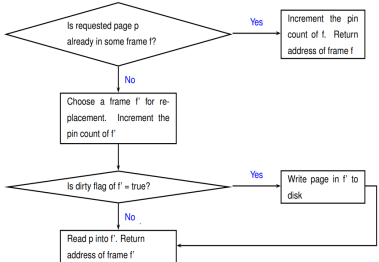
Magnetic Disks



- **Disk Access Time** Seek time + Rotational Latency + Transfer time
- **Response time** Queueing delay + Disk access time
- **Rotational Delay** $\frac{1}{2} \frac{60s}{RPM}$
- **Transfer Time** sectors on the same track * $\frac{TimePerRevolution}{SectorsPerTrack}$

Buffer Manager

- **Buffer pool** Main memory allocated for DBMS
- **pin count** is incremented upon pinning
- **dirty bit** is updated when the page is unpinned (if modified)
- Replacement is only possible if pin count == 0




Replacement Policies

LRU Policy

- Maintains a queue of pointers to frames with pin count = 0

Clock Replacement Policy

N = number of frames in buffer pool



CS3223, Sem 2, 2022/23

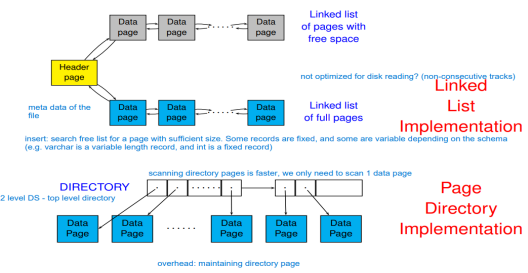
Buffer Manager

- Simplifies LRU with a second chance round robin system
- Each frame has a **reference bit** that is turned on when pin count reaches 0
- Replaces a page when referenced bit is off and pin count is 0

File Organisation


Heap File Implementations

worst case: O(n), and we need to scan through the entire data page to find out the space available



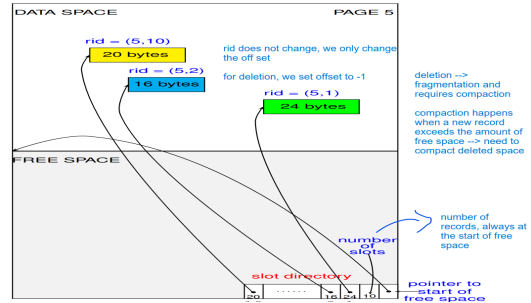
Page Formats: Fixed Length Records

- **Packed Organisation** Store records in contiguous slots
- **Unpacked Organisation** Uses a bit array to maintain free slots



Page Formats: Variable Length Records

- **Slotted Page** Store records in slots of record offset, record length
- **Record Offset:** Offset of the record from the start of the page



Record Formats

- **Fixed-Length Records**
 - Fields are stored consecutively
- **Variable-Length Records**
 - Delimit fields with special symbols
 - Use an array of field offsets

Each α_i is an offset to beginning of field F_i