

01. COMPUTATIONAL MODELS

- algorithm** → a well-defined procedure for finding the correct solution to the input
- correctness**
 - worst-case correctness** → correct on *every valid input*
 - other types of correctness: correct on random input/with high probability/approximately correct
- efficiency / running time** → measures the number of steps executed by an algorithm as a function of the *input size* (depends on computational model used)
 - number input: typically the length of binary representation
 - worst-case** running time → *max* number of steps executed when run on an input of size n

adversary argument →

inputs are decided such that they have different solutions

Comparison Model

- algorithm can **compare** any two elements in one time unit ($x > y, x < y, x = y$)
- running time = number of pairwise comparisons made
- array can be manipulated at no cost

Decision Tree

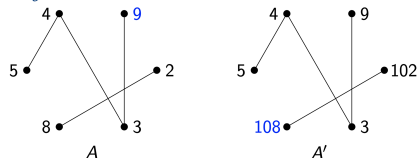
- each comparison represents the relationship between two elements
- each node is a comparison
- each branch is an outcome of the comparison
- log base is determined by the number of branches per node
- each leaf is a class label (decision after *all* comparisons)
- lower bound of worst-case** runtime = height of tree
- # of leaves = # of permutations → $\lg(n!) = \Theta(n \lg n)$
- any decision tree that can sort n elements must have height $\Omega(n \lg n)$.

Max Problem

problem: find largest element in array A of n distinct elements

Proof. $n - 1$ comparisons are needed

fix an algorithm M that solves the Max problem on all inputs using $< n - 1$ comparisons. construct graph G where nodes i and j are adjacent iff M compares i & j .




M cannot differentiate A and A' .

Second Largest Problem

problem: find the second largest element in $< 2n - 3$ comparisons ($2 \times \text{Maximum} \Rightarrow (n-1) + ((n-1)-1) = 2n-3$)

- solution: knockout tournament** → $n + \lceil \lg n \rceil - 2$



- bracket system: $n - 1$ matches
 - every non-winner has lost exactly once
- then compare the elements that have lost to the largest
 - the 2nd largest element must have lost to the winner
 - compares $\lceil \lg n \rceil$ elements that have lost to the winner using $\lceil \lg n \rceil - 1$ comparisons

Sorting

Claim. there is a sorting algorithm that requires $\leq n \lg n - n + 1$ comparisons.

Proof. every sorting algorithm must make $\geq \lg(n!)$ comparisons.

- let set \mathcal{U} be the set of all permutations of the set $\{1, \dots, n\}$ that the adversary could choose as array A . $|\mathcal{U}| = n!$
- for each query "is $A_i > A_j$?", if $\mathcal{U}_{yes} = \{A \in \mathcal{U} : A_i > A_j\}$ is of size $\geq |\mathcal{U}|/2$, set $\mathcal{U} := \mathcal{U}_{yes}$. else: $\mathcal{U} := \mathcal{U} \setminus \mathcal{U}_{yes}$
- the size of \mathcal{U} decreases by at most half with each comparison
- with $< \lg(n!)$ comparisons, \mathcal{U} will still contain at least 2 permutations

$$\begin{aligned} n! &\geq \left(\frac{n}{e}\right)^n \\ \Rightarrow \lg(n!) &\geq n \lg\left(\frac{n}{e}\right) = n \lg n - n \lg e \\ &\approx n \lg n - 1.44n \end{aligned}$$

⇒ roughly $n \lg n$ comparisons are **required** and **sufficient** for sorting n numbers

String Model

input	string of n bits
each query	find out one bit of the string

- n queries are **necessary** and **sufficient** to check if the input string is all 0s.
- query complexity** → number of bits of the input string queried by the algorithm
- evasive** → a problem requiring n query complexity

Graph Model

input	(symmetric) adjacency matrix of an n -node undirected graph
each query	find out if an edge is present between two chosen nodes (one entry of G)

- evasive** → requires $\binom{n}{2}$ queries

- Proof.* determining whether the graph is connected is evasive (requires $\binom{n}{2}$ queries)
 - suppose M is an algorithm making $\leq \binom{n}{2}$ queries.
 - whenever M makes a query, the algorithm tries not adding this edge, but adding all remaining unqueried edges.
 - if the resulting graph is connected, M replies 0 (i.e. edge does not exist)
 - else: replies 1 (edge exists)
 - after $< \binom{n}{2}$ queries, at least one entry of the adjacency matrix is unqueried.

02. ASYMPTOTIC ANALYSIS

- algorithm** → a *finite* sequence of well-defined instructions to solve a given computational problem
- word-RAM model** → runtime is the total number of instructions executed
 - operators, comparisons, if, return, etc
 - each instruction operates on a *word* of data (limited size) ⇒ fixed constant amount of time

Asymptotic Notations

upper bound (\leq): $f(n) = O(g(n))$

if $\exists c > 0, n_0 > 0$ such that $\forall n \geq n_0$,
 $0 \leq f(n) \leq cg(n)$

lower bound (\geq): $f(n) = \Omega(g(n))$

if $\exists c > 0, n_0 > 0$ such that $\forall n \geq n_0$,
 $0 \leq cg(n) \leq f(n)$

tight bound: $f(n) = \Theta(g(n))$

if $\exists c_1, c_2, n_0 > 0$ such that $\forall n \geq n_0$,
 $0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n)$

o -notation ($<$): $f(n) = o(g(n))$

if $\forall c > 0, \exists n_0 > 0$ such that $\forall n \geq n_0$,
 $0 \leq f(n) < cg(n)$

ω -notation ($>$): $f(n) = \omega(g(n))$

if $\forall c > 0, \exists n_0 > 0$ such that $\forall n \geq n_0$,
 $0 \leq cg(n) < f(n)$

Proof. $(n+1)! \neq O(n!)$ since $\frac{(n+1)!}{n!} = (n+1) > c$

Limits

Assume $f(n), g(n) > 0$.

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0 \quad \Rightarrow f(n) = o(g(n))$$

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} < \infty \quad \Rightarrow f(n) = O(g(n))$$

$$0 < \lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} < \infty \quad \Rightarrow f(n) = \Theta(g(n))$$

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} > 0 \quad \Rightarrow f(n) = \Omega(g(n))$$

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = \infty \quad \Rightarrow f(n) = \omega(g(n))$$

- Proof.*
- Since $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0$, we have for all $\epsilon > 0$, there exists $\delta > 0$ s.t. $\frac{f(n)}{g(n)} < \epsilon$ for $n > \delta$
 - Set $c = \epsilon$ and $n_0 = \delta$
 - $\forall n \geq n_0, \frac{f(n)}{g(n)} < c$
 - $\forall n \geq n_0, f(n) < cg(n)$
 - By definition, $f(n) = o(g(n))$

Properties of Big O

$$\Theta(g(n)) = O(g(n)) \cap \Omega(g(n))$$

- transitivity** - applies for $O, \Omega, \Theta, o, \omega$
 $f(n) = O(g(n)) \wedge g(n) = O(h(n)) \Rightarrow f(n) = O(h(n))$
- reflexivity** - for $O, \Omega, \Theta, f(n) = O(f(n))$
- symmetry** - $f(n) = \Theta(g(n)) \iff g(n) = \Theta(f(n))$
- complementarity** -
 - $f(n) = O(g(n)) \iff g(n) = \Omega(f(n))$
 - $f(n) = o(g(n)) \iff g(n) = \omega(f(n))$
- misc**
 - if $f(n) = \omega(g(n))$, then $f(n) = \Omega(g(n))$
 - if $f(n) = o(g(n))$, then $f(n) = O(g(n))$

$$\log \log n < \log n < (\log n)^k < n^k < (n+1)! < k^n$$

insertion sort: $O(n^2)$ with worst case $\Theta(n^2)$

03. ITERATION, RECURSION, DIVIDE-AND-CONQUER

Iterative Algorithms

- iterative** → loop(s), sequentially processing input elements
- loop invariant** implies correctness if
 - initialisation* - true before the first iteration of the loop
 - maintenance* - if true before an iteration, it remains true at the beginning of the next iteration
 - termination* - true when the algorithm terminates

examples

- insertionSort:** with loop variable as $j, A[1..j-1]$ is sorted.
 - $A[1..j] = A[1..j]$. Elements not considered are unaffected.
 - $A[i+2..j] = A[i+1..j-1]$. Relative order of shifted elements is preserved.
 - $A[i+2..j] > \text{key}$. Elements to its right are sorted and greater.
- selectionSort:** with loop variable as j , the array $A[1..j-1]$ is sorted and contains the $j-1$ smallest elements of A .
- Dijkstra's:**

Proof.

 - invariant 1** $\forall x \in R: \text{dist}[x] = \sigma(s, x)$
 - invariant 2** $\forall y$ neighbouring $x \in R$:
 $\text{dist}[y] = \min_{x \in R} \sigma(s, x) + W(x, y)$

Recursive Algorithm

- recursive** → solves sub problems
- Correctness is proven using **mathematical induction** on size of problem
- Use strong induction, prove base case, show algorithm works assuming it works for all smaller cases

Examples

- BINARY-SEARCH** (A, a, b, x) $\triangleright A[a \dots b]$
if $a > b$ then
 return *false*
else $mid = \lfloor (a+b)/2 \rfloor$
if $x == A[mid]$ then
 return *true*
if $x < A[mid]$ then
 return **BINARY-SEARCH** ($A, a, mid-1, x$)
else
 return **BINARY-SEARCH** ($A, mid+1, b, x$)
- **binary search**(A,a,b,x) returns the correct answer when b-a+1=n
 - Base case: n=b-a+1=0, since a=b+1, A[a..b] is empty and the answer is false
 - Inductive step: $n = b - a + 1 > 0$
 - By strong induction, assume (A,a',b',x) returns the correct answer for all j s.t $0 \leq j \leq n-1$ where $j = b' - a' + 1$
 - By the algorithm, $mid = \lfloor \frac{a+b}{2} \rfloor$ and $a \leq mid \leq b$
 - If $x == A[mid]$ then $x \in A[a..b]$ and the algorithm returns true correctly
 - If $x < A[mid]$ then $x \in A[a..mid-1]$ if $x \in A[a..b]$
 - By the inductive hypothesis, (A,a,mid-1,x) is correct since $0 \leq (mid-1) - a + 1 = mid - a \leq n-1$
 - The $x > A[mid]$ is similar

Divide-and-Conquer
powering a number

- problem:* compute $f(n, m) = a^n \pmod m$ for all $n, m \in \mathbb{Z}$
- observation: $f(x+y, m) = f(x, m) * f(y, m) \pmod m$
 - **naive solution:** recursively compute and combine $f(n-1, m) * f(1, m) \pmod m$
 - $T(n) = T(n-1) + T(1) + \Theta(1) \Rightarrow T(n) = \Theta(n)$
 - **better solution:** divide and conquer (only one sub problem computed)
 - divide: trivial
 - conquer: recursively compute $f(\lfloor n/2 \rfloor, m)$
 - combine:
 - $f(n, m) = f(\lfloor n/2 \rfloor, m)^2 \pmod m$ if n is even
 - $f(n, m) = f(1, m) * f(\lfloor n/2 \rfloor, m)^2 \pmod m$ if odd
 - $T(n) = T(n/2) + \Theta(1) \Rightarrow \Theta(\log n)$

Peak finding

- problem:* Find the peak element (no neighbours are greater) in 2D array
- **Naive** O(mn)
 - for the middle column, find the maximum element
 - return if it is peak
 - p1=Find2DPeak(left)
 - p2=Find2DPeak(right)
 - return p1 or p2 if one is a peak
 - **Divide and conquer** O(mlogn), $T(n) = T(n/2) + O(1)$ n is number of columns
 - find the middle column, find the maximum element
 - if it is a peak, return it
 - if not, recurse on the side with a larger element
 - **Optimised** O(m+n)

- find the middle column, find the maximum element
- recurse on the quarter with the larger element

Solving Recurrences

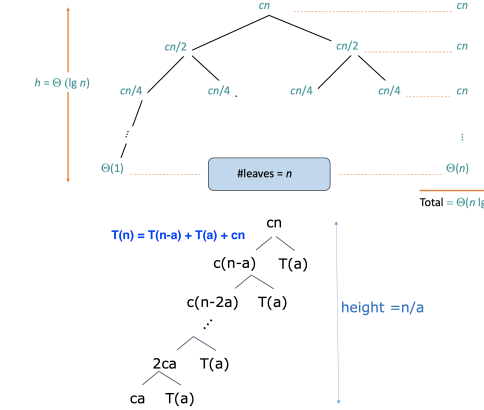
for a sub-problems of size $\frac{n}{b}$ where $f(n)$ is the time to divide and combine,
$$T(n) = aT(\frac{n}{b}) + f(n)$$

Telescoping

Express $\frac{T(n)}{g(n)}$ as $\frac{T(\frac{n}{b})}{g(\frac{n}{b})} + h(n)$

1. Example: $T(n) = 2T(n/2) + n$
2. $\frac{T(n)}{n} = \frac{T(n/2)}{n/2} + 1$
3. $\frac{T(n)}{n} = \frac{T(n/2)}{n/2} + 1$
 $\frac{T(n/2)}{n/2} = \frac{T(n/4)}{n/4} + 1$
 $\frac{T(n/4)}{n/4} = \frac{T(n/8)}{n/8} + 1$
 ...
 $\frac{T(2)}{2} = \frac{T(1)}{1} + 1$
Hence, $T(n) = O(n \log n)$.

- If $g(n)=n$, we need $a=b$ since $g(n) = n^{\log_b(a)}$
- Recursion tree**
total = height \times number of leaves
• each node represents the cost of a single subproblem
• height of the tree = longest path from root to leaf



Master method

- $a \geq 1, b > 1$, and f is asymptotically positive.
 - a (number of sub problems), b (size of sub problems), f (time to divide and combine)
 $T(n) = aT(\frac{n}{b}) + f(n) =$
 - $\Theta(n^{\log_b a})$ if $f(n) < n^{\log_b a}$ polynomially
 - $\Theta(n^{\log_b a} \log n)$ if $f(n) = n^{\log_b a}$
 - $\Theta(f(n))$ if $f(n) > n^{\log_b a}$ polynomially
- three common cases**
1. If $f(n) = O(n^{\log_b a - \epsilon})$ for some constant $\epsilon > 0$,
 - $f(n)$ grows polynomially slower than $n^{\log_b a}$ by n^ϵ factor.
 - then $T(n) = \Theta(n^{\log_b a})$.
 - This is when overhead at leaf $>$ overhead at root

2. If $f(n) = \Theta(n^{\log_b a} \log^k n)$ for some $k \geq 0$,
 - $f(n)$ and $n^{\log_b a}$ grow at similar rates.
 - then $T(n) = \Theta(n^{\log_b a} \log^{k+1} n)$
3. If $f(n) = \Omega(n^{\log_b a + \epsilon})$ for some constant $\epsilon > 0$,
 - and $f(n)$ satisfies the **regularity condition**
 - $af(n/b) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n
 - this guarantees that the sum of subproblems is smaller than $f(n)$.
 - $f(n)$ grows polynomially faster than $n^{\log_b a}$ by n^ϵ factor
 - then $T(n) = \Theta(f(n))$.
 - This is when the root (splitting) $>$ leaf

Substitution method

1. guess that $T(n) = O(f(n))$.
2. verify by induction:
2.1. to show that for $n \geq n_0, T(n) \leq c \cdot f(n)$
2.2. set $c = \max\{2, n_0\}$ and $n_0 = 1$
2.3. verify base case(s): $T(n_0) = q$
2.4. recursive case ($n > n_0$):
 - by strong induction, assume $T(k) \leq c \cdot f(k)$ for $n > k \geq n_0$
 - $T(n) = \text{recurrence}_T \dots \leq c \cdot f(n)$
- 2.5. hence $T(n) = O(f(n))$.

! may not be a tight bound!

example

Proof. $T(n) = 4T(n/2) + n^2 / \lg n \Rightarrow \Theta(n^2 \lg \lg n)$

$$\begin{aligned} T(n) &= 4T(n/2) + \frac{n^2}{\lg n} \\ &= 4(4T(n/4) + \frac{(n/2)^2}{\lg n - \lg 2}) + \frac{n^2}{\lg n} \\ &= 16T(n/4) + \frac{n^2}{\lg n - \lg 2} + \frac{n^2}{\lg n} \\ &= \sum_{k=1}^{\lg n} \frac{n^2}{\lg n - k} \\ &= n^2 \lg \lg n \text{ by approx. of harmonic series } (\sum \frac{1}{k}) \end{aligned}$$

Can also be solved via telescoping using $g(n)=n^2$

Proof. $T(n) = 4T(n/2) + n \Rightarrow O(n^2)$

To show that for all $n \geq n_0, T(n) \leq c_1 n^2 - c_2 n$

1. Set $c_1 = q + 1, c_2 = 1, n_0 = 1$.
2. Base case ($n = 1$): subbing into $c_1 n^2 - c_2 n$,
 $T(1) = q \leq (q+1)(1)^2 - (1)(1)$
3. Recursive case ($n > 1$):
- by strong induction, assume $T(k) \leq c_1 \cdot k^2 - c_2 \cdot k$ for all $n > k \geq 1$
 - $T(n) = 4T(n/2) + n$
 - $= 4(c_1(n/2)^2 - c_2(n/2)) + n$
 - $= c_1 n^2 - 2c_2 n + n$
 - $= c_1 n^2 - c_2 n + (1 - c_2)n$
 - $= c_1 n^2 - c_2 n$ since $c_2 = 1 \Rightarrow 1 - c_2 = 0$

Non-comparison sort

Counting sort

Set $C[i]$ be the number of elements equal i .

Set $C[i]$ be the number of elements smaller than or equal i .

Move elements equal i to $B[C[i-1]+1..C[i]]$.

for $i \leftarrow 1$ to k
 do $C[i] \leftarrow 0$
for $j \leftarrow 1$ to n
 do $C[A[j]] \leftarrow C[A[j]] + 1 \triangleright C[i] = |\{\text{key} = i\}|$
for $i \leftarrow 2$ to k *partial sum*
 do $C[i] \leftarrow C[i] + C[i-1] \triangleright C[i] = |\{\text{key} \leq i\}|$
for $j \leftarrow n$ downto 1
 do $B[C[A[j]]] \leftarrow A[j]$
 do $C[A[j]] \leftarrow C[A[j]] - 1$

- $O(n+k)$, where k is the number of elements
- Stable

Radix sort

Suppose there are T digits.

3	2	9	7	2	0	7	2	0	3	2	9
4	5	7	3	5	5	3	2	9	3	5	5
6	5	7	4	3	6	4	3	6	4	3	6
8	3	9	4	5	7	8	3	9	4	5	7
4	3	6	6	5	7	3	5	5	6	5	7
7	2	0	3	2	9	4	5	7	7	2	0
3	5	5	8	3	9	6	5	7	8	3	9

- Example:** We can use counting sort as the stable sorting algorithm.
- Each pass processes r digits at $O(n + 2^r)$
 - $T(n, b) = \theta(\frac{b}{2} (n + 2^r))$ where b is the number of bits in the key
 - Optimal $r = \lg n$ $T(n, b) = \theta(\frac{bn}{\lg n})$
 - Fast when the number of passes ($\frac{b}{r}$) is small
 - Little locality of reference, not cache friendly

04. AVERAGE-CASE ANALYSIS & RANDOMISED ALGORITHMS

- **average case** $A(n) \rightarrow$ expected running time when the input is chosen uniformly at random from the set of all $n!$ permutations
 - $A(n) = \frac{1}{n!} \sum_{\pi} Q(\pi)$ where $Q(\pi)$ is the time complexity when the input is permutation π .
 - $A(n) = \mathbb{E}_{x \sim \mathcal{D}_n} [\text{Runtime of Alg on } x]$
 - $\mathbb{E}_{x \sim \mathcal{D}_n}$ is a probability distribution on U restricted to inputs of size n .

Quicksort Analysis

- divide & conquer, linear-time $\Theta(n)$ partitioning subroutine
- assume we select the first array element as pivot
- $T(n) = T(j) + T(n-j-1) + \Theta(n)$
 - if the pivot produces subarrays of size j and $(n-j-1)$
- **worst-case:** $T(n) = T(0) + T(n-1) + \Theta(n) \Rightarrow \Theta(n^2)$

Proof. for quicksort, $A(n) = O(n \log n)$

let $P(i)$ be the set of all those permutations of elements $\{e_1, e_2, \dots, e_n\}$ that begins with e_i .

Let $G(n, i)$ be the average running time of quicksort over $P(i)$. Then
 $G(n) = A(i-1) + A(n-i) + (n-1)$
 $A(n) = \frac{1}{n} \sum_{i=1}^n G(n, i)$
 $= \frac{1}{n} \sum_{i=1}^n (A(i-1) + A(n-i) + (n-1))$
 $= \frac{2}{n} \sum_{i=1}^n A(i-1) + n-1$

= $O(n \log n)$ by taking it as area under integration

quicksort vs mergesort

	average	best	worst
quicksort	$1.39n \lg n$	$n \lg n$	$n(n-1)$
mergesort	$n \lg n$	$n \lg n$	$n \lg n$

- disadvantages of mergesort:
 - overhead of temporary storage
 - cache misses
- advantages of quicksort
 - in place
 - reliable (as $n \uparrow$, chances of deviation from avg case \downarrow)
- issues with quicksort
 - distribution-sensitive** \rightarrow time taken depends on the initial (input) permutation. Resolved with median pivot or randomised partitions

Randomised Algorithms

- randomised algorithms** \rightarrow output and running time are **functions** of the **input** and **random bits chosen**
 - vs non-randomised: output & running time are functions of the *input only*
- expected running time = worst-case running time = $E(n) = \max_{\text{input } x \text{ of size } n} \mathbb{E}[\text{Runtime of RandAlg on } x]$
- randomised quicksort**: choose pivot at random
 - probability that the runtime of *randomised* quicksort exceeds average by $x\%$ = $n^{-\frac{x}{100} \ln \ln n}$
 - P(time takes at least double of the average) = 10^{-15}
 - distribution insensitive

Balls into bins - Indicator Random Variable

There are n balls and m bins. Each ball is placed into a bin at random. How many empty bins?

- $X_i = 1$ if ball i is in bin j , 0 otherwise
- $E(X_i) = 1 * P(i^{th} \text{ bin empty}) + 0 * P(\dots) = (1 - \frac{1}{n})^m$
- $E(X) = E(X_1) + E(X_2) + \dots + E(X_n) = n(1 - \frac{1}{n})^m$

Randomised Quicksort Analysis

$T(n) = n - 1 + T(q - 1) + T(n - q)$
Let $A(n) = \mathbb{E}[T(n)]$ where the expectation is over the randomness in expectation.
Taking expectations and applying linearity of expectation:
 $A(n) = n - 1 + \frac{1}{n} \sum_{q=1}^n (A(q - 1) + A(n - q))$
 $= n - 1 + \frac{2}{n} \sum_{q=1}^{n-1} A(q)$

$A(n) = n \log n \Rightarrow$ same as average case quicksort

Randomised Quickselect

- $O(n)$ to find the k^{th} smallest element
- randomisation: unlikely to keep getting a bad split

Types of Randomised Algorithms

- randomised **Las Vegas** algorithms
 - output is always correct
 - runtime is a *random variable*
 - e.g. randomised quicksort, randomised quickselect
- randomised **Monte Carlo** algorithms

- output may be incorrect with some small probability
- runtime is *deterministic*

Examples

- smallest enclosing circle**: given n points in a plane, compute the smallest radius circle that encloses all n points
 - best **deterministic** algorithm: $O(n)$, but complex
 - Las Vegas: average $O(n)$, simple solution
- minimum cut**: given a connected graph G with n vertices and m edges, compute the smallest set of edges whose removal would disconnect G .
 - best **deterministic** algorithm: $O(mn)$
 - Monte Carlo**: $O(m \log n)$, error probability n^{-c} for any c
- primality testing**: determine if an n bit integer is prime
 - best **deterministic** algorithm: $O(n^6)$
 - Monte Carlo**: $O(kn^2)$, error probability 2^{-k} for k checks

Geometric Distribution

Let X be the number of trials repeated until success. X is a random variable and follows a geometric distribution with probability p .

Expected number of trials, $E[X] = \frac{1}{p}$
 $Pr[X = k] = q^{k-1}p$

Linearity of Expectation

For any two events X, Y and a constant a ,
 $E[X + Y] = E[X] + E[Y]$
 $E[aX] = aE[X]$

Coupon Collector Problem

n types of coupon are put into a box and randomly drawn with replacement. What is the expected number of draws needed to collect at least one of each type of coupon?

- let T_i be the time to collect the i -th coupon after the $i - 1$ coupon has been collected.
 - Probability of collecting a new coupon, $p_i = \frac{(n-(i-1))}{n}$
 - T_i has a **geometric distribution**
 - $E[T_i] = 1/p_i$

total number of draws, $T = \sum_{i=1}^n T_i$

$E[T] = E[\sum_{i=1}^n T_i] = \sum_{i=1}^n E[T_i]$ by linearity of expectation

$= \sum_{i=1}^n \frac{n}{n-(i-1)} = n \cdot \sum_{i=1}^n \frac{1}{i} = \Theta(n \log n)$

05. HASHING

Dictionary ADT

- different types:
 - static** - fixed set of inserted items; only care about queries
 - insertion-only** - only insertions and queries
 - dynamic** - insertions, deletions, queries
- implementations
 - sorted list (static) - $O(\log N)$ query
 - balanced search tree (dynamic) - $O(\log N)$ all operations
 - direct access table

- \times needs items to be represented as non-negative integers (**prehashing**)
- \times huge space requirement
- using \mathcal{H} for dictionaries: need to store both the hash table and the matrix A .
 - additional storage overhead = $\Theta(\log N \cdot \log |U|)$, if $M = \Theta(N)$
 - other universal hashing constructions may have more efficient hash function evaluation
- associative array** - has both key and value (dictionary in this context has only key)

Hashing

- hash function**, $h : U \rightarrow \{1, \dots, M\}$ gives the location of where to store in the hash table
 - notation: $[M] = \{1, \dots, M\}$ $[M] = \{1, \dots, M\}$
 - storing N items in hash table of size M
- collision** \rightarrow for two different keys x and y , $h(x) = h(y)$
 - resolve by **chaining**, **open addressing**, etc
- desired properties
 - \checkmark minimise collisions - `query(x)` and `delete(x)` take time $\Theta(|h(x)|)$
 - \checkmark minimise storage space - aim to have $M = O(N)$
 - \checkmark function h is easy to compute (assume constant time)
- if $|U| \geq (N - 1)M + 1$, for any $h : U \rightarrow [M]$, there is a set of N elements having the same hash value.
 - Proof**: pigeonhole principle
- use **randomisation** to overcome the adversary
 - e.g. randomly choose between two *deterministic* hash functions h_1 and h_2
 \Rightarrow for any pair of keys, with probability $\geq \frac{1}{2}$, there will be no collision

Universal Hashing

Suppose \mathcal{H} is a set of hash functions mapping U to $[M]$.
 \mathcal{H} is **universal** if $\forall x \neq y, \frac{|h \in \mathcal{H} : h(x) = h(y)|}{|\mathcal{H}|} \leq \frac{1}{M}$
or $Pr_{h \sim \mathcal{H}} [h(x) = h(y)] \leq \frac{1}{M}$

- aka: for any $x \neq y$, if h is chosen uniformly at random from a universal \mathcal{H} , then there is at most $\frac{1}{M}$ probability that $h(x) = h(y)$
- probability where h is sampled uniformly from \mathcal{H}
- aka: for any $x \neq y$, the fraction of hash functions with collisions is at most $\frac{1}{M}$.

Properties of universal hashing

Collision Analysis

- for any N elements $x_1, \dots, x_N \in U$, the **expected number of collisions** between x_N and other elements is $< N/M$.
 - it follows that for K operations, the expected cost of the last operation is $< K/M = O(1)$ if $M > K$.

Proof. by definition of Universal Hashing, each element $x_1, \dots, x_{N-1} \in U$ has at most $\frac{1}{M}$ probability of collision with x_N (over random choice of h).
by indicator r.v., $E[A_i] = P(A_i = 1) \leq \frac{1}{M}$. expected number of collisions = $(N - 1) \cdot \frac{1}{M} < \frac{N}{M}$.

- if x_1, \dots, x_N are added to the hash table, and $M > N$, the expected **number of pairs** (i, j) with collisions is $< 2N$.

Proof. let A_{ij} be an indicator r.v. for collision.

$$\mathbb{E}[\sum_{1 \leq i, j \leq N} A_{ij}] = \sum_{i=1}^N \mathbb{E}[A_{ii}] + \sum_{i \neq j} \mathbb{E}[A_{ij}] \leq N \cdot 1 + N(N - 1) \cdot \frac{1}{M} < 2N$$

Expected Cost

- for any sequence of N operations, if $M > N$, then the **expected total cost** for executing the sequence is $O(N)$.

Proof. linearity of expectation; sum up expected costs

Construction of Universal Family

Obtain a universal family of hash functions with $M = O(N)$.

- Suppose U is indexed by u -bit strings and $M = 2^m$.
- For any $m \times u$ binary matrix A , $h_A(x) = Ax \pmod{2}$
 - each element $x \Rightarrow x \% 2$
 - x is a $u \times 1$ matrix $\Rightarrow Ax$ is $m \times 1$
- Claim**: $\{h_A : A \in \{0, 1\}^{m \times u}\}$ is universal
- e.g. $U = \{00, 01, 10, 11\}$, $M = 2$
 - h_{ab} means $A = \begin{bmatrix} a & b \end{bmatrix}$

	00	01	10	11
h_{00}	0	0	0	0
h_{01}	0	1	0	1
h_{10}	0	0	1	1
h_{11}	0	1	1	0

Proof. Let $x \neq y$. Let $z = x - y$. We know $z \neq 0$.

Collision: $P(Ax = Ay) = P[A(x - y) = 0] = P(Az = 0)$.

To show $P(Az = 0) \leq \frac{1}{M}$.

Special case - Suppose z is 1 at the i -th coordinate but 0 everywhere else. Then Az is the i -th column of A . Since the i -th column is uniformly random, $P(Az = 0) = \frac{1}{2^m} = \frac{1}{M}$.

General case - Suppose z is 1 at the i -th coordinate. Let $z = [z_1 \ z_2 \ \dots \ z_u]^T$. $A = [A_1 \ A_2 \ \dots \ A_u]$ hence Az_k is the k -th column of A .
Then $Az = z_1 A_1 + z_2 A_2 + \dots + z_u A_u$.
 $Az = 0 \Rightarrow z_1 A_1 = -(z_2 A_2 + \dots + z_u A_u)$ (*)
We fix $z_1 A_1$ to be an arbitrary $m \times 1$ matrix of 1s and 0s. The probability that (*) holds is $\frac{1}{2^m}$.

Perfect Hashing

static case - N fixed items in the dictionary x_1, x_2, \dots, x_N
To perform Query in $O(1)$ worst-case time.

Quadratic Space: $M = N^2$

if \mathcal{H} is universal and $M = N^2$, and h is sampled uniformly from \mathcal{H} , then the expected number of collisions is < 1 .

Proof. for $i \neq j$, let indicator r.v. A_{ij} be equal to 1 if $h(x_i) = h(x_j)$, or 0 otherwise.
By universality, $E[A_{ij}] = P(A_{ij} = 1) \leq 1/N^2$
 $E[\# \text{ collisions}] = \sum_{i < j} E[A_{ij}] \leq \binom{N}{2} \frac{1}{N^2} < 1$

It follows that there exists $h \in \mathcal{H}$ causing no collisions (because if not, $\mathbb{E}[\# \text{ collisions}]$ would be ≥ 1).

2-Level Scheme: $M = N$

- No collision and less space needed

Construction

- Choose $h : U \rightarrow [N]$ from a universal hash family.
- Let L_k be the number of x_i 's for which $h(x_i) = k$.
 - Choose h_1, \dots, h_N **second-level** hash functions $h_k : [N] \rightarrow [(L_k)^2]$ s.t. there are no collisions among the L_k elements mapped to k by h .
 - quadratic second-level table \rightarrow ensures no collisions using quadratic space

Analysis

if \mathcal{H} is universal and h is sampled uniformly from \mathcal{H} , then

$$E\left[\sum_k L_k^2\right] < 2N$$

Proof. For $i, j \in [1, N]$, define indicator r.v. $A_{ij} = 1$ if $h(x_i) = h(x_j)$, or 0 otherwise.

$$A_{ij} = \# \text{ possible collisions} = \# \text{ pairs} \cdot 2 = L_k^2$$

$$\text{Hence } \sum_k L_k^2 = \sum_{i,j} A_{ij}$$

$$\begin{aligned} E[\sum_{i,j} A_{ij}] &= \sum_i E[A_{ii}] + \sum_{i \neq j} E[A_{ij}] \\ &\leq N \cdot 1 + N(N-1) \cdot \frac{1}{N} \\ &< 2N \end{aligned}$$

Hash Table Resizing

- when number of inserted items, N is not known
 - rehashing** - choose a new hash function of a larger size and re-hash all elements
 - costly but infrequent \Rightarrow amortize

06. FINGERPRINTING & STREAMING

String Pattern Matching

problem: does the pattern string P occur as a substring of the text string T ?

- m = length of P , n = length of T , ℓ = size of alphabet
- assumption: operations on strings of length $O(\log n)$ can be executed in $O(1)$ time. (word-RAM model)
- naive solution: $\Theta(n^2)$

Fingerprinting approach (Karp-Rabin)

- faster string equality check:
 - for substring X , check $h(X) == h(P)$ for a hash function $h \Rightarrow \Theta(1)$ + cost of hashing instead of $\Theta(|X|)$
- Rolling Hash:** $O(m + n)$
 - update the hash from what we already have from the previous hash - $O(1)$
 - compute $n - m + 1$ hashes in $O(n)$ time
 - Monte Carlo algorithm

Division Hash

Choose a random **prime** number p in the range $\{1, \dots, K\}$.
For integer x , $h_p(x) = x \pmod p$

- if p is small and x is b -bits long in binary, hashing $\Rightarrow O(b)$
- hash family $\{h_p\}$ is approximately universal
- if $0 \leq x < y < 2^b$, then $Pr_h[h_p(x) = h_p(y)] < \frac{b \ln K}{K}$

Proof. $h_p(x) = h_p(y)$ when $y - x = 0 \pmod p$.

$$\text{Let } z = y - x.$$

Since $z < 2^b$, then z can have at most b distinct prime factors.

p divides z if p is one of these $\leq b$ prime factors.

number of primes in range $\{1, \dots, K\}$ is $> \frac{K}{\ln K}$,

hence the probability is $b / \frac{K}{\ln K} = \frac{b \ln K}{K}$

values of K

- higher K = lower probability of false positive
 - for $\delta = \frac{1}{100n}$, $P(\text{false positive}) \leq 1\%$.

$$\forall \delta > 0, \text{ if } X \neq Y \text{ and } K = \frac{2m}{\delta} \cdot \lg \ell \cdot \lg(\frac{2m}{\delta} \lg \ell), \text{ then } Pr[h(X) = h(Y)] < \delta$$

Streaming

problem: Consider a sequence of insertions or deletions of items from a large universe \mathcal{U} . At the end of the stream, the frequency f_i of item i is its net count.

Let M be the sum of all frequencies at the end of stream.

naive solutions

- direct access table - $\Omega(U)$ space
- sorted list - $\Omega(M)$ space, no $O(1)$ update
- binary search tree - $O(M)$ space

Frequency Estimation

an approximation \hat{f}_i is **ϵ -approximate** if

$$f_i - \epsilon M \leq \hat{f}_i \leq f_i + \epsilon M$$

Using Hash Table

$$f_i \leq \mathbb{E}[\hat{f}_i] \leq f_i + M/k$$

- increment/decrement $A[h(j)]$ on an empty table A of size k
- collision \Rightarrow false positives \Rightarrow may give overestimate of f_i
 - $A[h(i)] = \sum_{j:h(j)=h(i)} f_j \geq f_i$
- if h is drawn from a universal family, overestimate, $\mathbb{E}[A[h(i)] - f_i] \leq M/k$
- space: $O(\frac{1}{\epsilon} \cdot \lg M + \lg U \cdot \lg M)$

let $k = \frac{1}{\epsilon}$ for some $\epsilon > 0$.

- number of rows = $O(\frac{1}{\epsilon})$
- size of each row = $O(\lg M)$
- size of hash function (using universal hash family from ch.05) = $O(\lg U \cdot \lg M)$

- Count-Min Sketch** \rightarrow gives a bound on the probability that \hat{f}_i deviates from f_i instead of a bound on the expectation of the gap

07. AMORTIZED ANALYSIS

- amortized analysis** \rightarrow guarantees the *average* performance of each operation in the *worst case*.
- total amortized cost provides an *upper bound* on the total true cost
- For a sequence of n operations o_1, o_2, \dots, o_n ,
 - let $t(i)$ be the time complexity of the i -th operation o_i
 - let $f(n)$ be the *worst-case* time complexity for *any* of the n operations
 - let $T(n)$ be the time complexity of all n operations

$$T(n) = \sum_{i=1}^n t(i) = nf(n)$$

Types of Amortized Analysis

Aggregate method

- look at the whole sequence, sum up the cost of operations and take the average - simpler but less precise
- e.g. binary counter - amortized $O(1)$
- e.g. queues (with INSERT and EMPTY) - amortized $O(1)$
- Find (a) The number of operations and (b) the upperbound of each operation
 - $a = n$
 - $b = \sum_{i=1}^n t(i) = nf(n)$

Accounting method

- charge the i -th operation a fictitious amortized cost $c(i)$
 - amortized cost** $c(i)$ is a fixed cost for each operation
 - true cost** $t(i)$ depends on when the operation is called
- amortized cost $c(i)$ must satisfy:

$$\sum_{i=1}^n t(i) \leq \sum_{i=1}^n c(i) \text{ for all } n$$

- take the extra amount for cheap operations early on as "credit" paid in advance for expensive operations
 - invariant:** bank balance never drops below 0
- the total amortized cost provides an **upper bound** on the total true cost

Potential method

- ϕ : potential function associated with the algo/DS
- $\phi(i)$: potential at the end of the i -th operation
- c_i : amortized cost of the i -th operation
- t_i : true cost of the i -th operation

$$\begin{aligned} c_i &= t_i + \phi(i) - \phi(i-1) \\ \sum_{i=1}^n c_i &= \phi(n) - \phi(0) + \sum_{i=1}^n t_i \end{aligned}$$


- hence as long as $\phi(n) \geq 0$, then amortized cost is an upper bound of the true cost.

$$\sum_{i=1}^n c_i \geq \sum_{i=1}^n t_i$$

- Validity** $\phi(0) = 0$ and $\phi(i) \geq 0$ for all i
- e.g.** for queue:
 - let $\phi(i)$ = # of elements in queue after the i -th operation
 - amortized cost for insert:
$$c_i = t_i + \phi(i) - \phi(i-1) = 1 + 1 = 2$$
 - amortized cost for empty (for k elements):
$$c_i = t_i + \phi(i) - \phi(i-1) = k + 0 - k = 0$$
- try to keep $c(i)$ small: using $c(i) = t(i) + \Delta\phi_i$
 - if $t(i)$ is small, we want $\Delta\phi_i$ to be positive and small
 - if $t(i)$ is large, we want $\Delta\phi_i$ to be negative and large

e.g. Dynamic Table (insertion only)

Aggregate method




Accounting method

- charge \$3 per insertion
 - \$1 for insertion itself
 - \$1 for moving itself when the table expands

- \$1 for moving one of the existing items when the table expands

Potential method



Operation Insert(·)	Actual Cost	$\Delta\phi$	Amortized Cost
Case 1: when table is not full	1	2	3
Case 2: when table is already full	i	$3 - i$	3

$$\begin{aligned} \text{Amortized cost of } n \text{ insertions} &= 3n = O(n) \\ \text{Actual cost of } n \text{ insertions} &= O(n) \end{aligned}$$

- show that SUM of amortized cost \geq SUM of actual cost
- conclude that sum of amortized cost is $O(f(n)) \Rightarrow$ sum of actual cost is $O(f(n))$

08. DYNAMIC PROGRAMMING

- cut-and-paste proof** \rightarrow proof by contradiction - suppose you have an optimal solution. Replacing ("cut") subproblem solutions with this subproblem solution ("paste" in) should improve the solution. If the solution doesn't improve, then it's not optimal (contradiction).
- overlapping subproblems** - recursive solution contains a small number of distinct subproblems repeated many times
- optimal substructure** - optimal solution to a problem contains optimal solutions to subproblems

Longest Common Subsequence

- for sequence $A : a_1, a_2, \dots, a_n$ stored in array
- C is a **subsequence** of $A \rightarrow$ if we can obtain C by removing zero or more elements from A .

problem: given two sequences $A[1..n]$ and $B[1..m]$, compute the *longest* sequence C such that C is a subsequence of A and B .

brute force solution

- check *all* possible subsequences of A to see if it is also a subsequence of B , then output the longest one.
- analysis: $O(m2^n)$
 - checking each subsequence takes $O(m)$
 - 2^n possible subsequences

recursive solution


let $LCS(i, j)$: longest common subsequence of $A[1..i]$ and $B[1..j]$

- base case: $LCS(i, 0) = \emptyset$ for all i , $LCS(0, j) = \emptyset$ for all j
- general case:
 - if last characters of A, B are $a_n = b_m$, then $LCS(n, m)$ must terminate with $a_n = b_m$
 - the optimal solution will match a_n with b_m
 - if $a_n \neq b_m$, then either a_n or b_m is not the last symbol
- optimal substructure:** (general case)
 - if $a_n = b_m$,
$$LCS(n, m) = LCS(n-1, m-1) :: a_n$$
 - if $a_n \neq b_m$,
$$LCS(n, m) = LCS(n-1, m) \parallel LCS(n, m-1)$$

- simplified problem:**
 - $L(n, m) = 0$ if $n = 0$ or $m = 0$
 - if $a_n = b_m$, then $L(n, m) = L(n - 1, m - 1) + 1$
 - if $a_n \neq b_m$, then $L(n, m) = \max(L(n, m - 1), L(n - 1, m))$
- analysis**
 - number of distinct subproblems = $(n + 1) \times (m + 1)$
 - to use $O(\min\{m, n\})$ space: bottom-up approach, column by column
 - memoize for DP \Rightarrow makes it $O(mn)$ instead of exponential time

Knapsack Problem

- input: $(w_1, v_1), (w_2, v_2), \dots, (w_n, v_n)$ and capacity W
- output: subset $S \subseteq \{1, 2, \dots, n\}$ that maximises $\sum_{i \in S} v_i$ such that $\sum_{i \in S} w_i \leq W$



- 2^n subsets \Rightarrow naive algorithm is costly
- recursive solution:**
 - let $m[i, j]$ be the maximum value that can be obtained using a subset of items $\{1, 2, \dots, i\}$ with total weight no more than j .
 - $m[i, j] = \begin{cases} 0, & \text{if } i = 0 \text{ or } j = 0 \\ \max\{m[i-1, j-w_i] + v_i, m[i-1, j]\}, & \text{if } w_i \leq j \\ m[i-1, j], & \text{otherwise} \end{cases}$
- analysis:** $O(nW)$
 - $O(nW)$ is **not** a polynomial time algorithm
 - not polynomial in input bitsize
 - W can be represented in $O(\lg W)$ bits
 - n can be represented in $O(\lg n)$ bits
 - polynomial time is strictly in terms of the number of bits for the input

Changing Coins

problem: use the fewest number of coins to make up n cents using denominations d_1, d_2, \dots, d_n . Let $M[j]$ be the fewest number of coins needed to change j cents.

- optimal substructure:**

$$M[j] = \begin{cases} 1 + \min_{i \in [k]} M[j - d_i], & j > 0 \\ 0, & j = 0 \\ \infty, & j < 0 \end{cases}$$

Proof. Suppose $M[j] = t$, meaning $j = d_{i_1} + d_{i_2} + \dots + d_{i_t}$ for some $i_1, \dots, i_t \in \{1, \dots, k\}$. Then, if $j' = d_{i_1} + d_{i_2} + \dots + d_{i_{t-1}}$, $M[j'] = t - 1$, because otherwise if $M[j'] < t - 1$, by **cut-and-paste** argument, $M[j] < t$.

- runtime: $O(nk)$ for n cents, k denominations

Dijkstra's Algorithm

- property 1:** The nearest neighbor is also the vertex nearest to s
- Violated by **negative edges** since there can exist a shorter path to a vertex that is not the nearest neighbor
- property 2:** Optimal subpath property
- Violated when there are **negative cycles** since the final parth can consist of a subparth that is not the optimal subpath

09. GREEDY ALGORITHMS

- solve only one subproblem at each step
- beats DP and divide-and-conquer when it works
- greedy-choice property** \rightarrow a locally optimal choice is globally optimal
- Note:** This is not true for DP, a clear counter-example is that the LCS of a substring is not necessarily in the LCS of the entire string

Examples

Fractional Knapsack

- $O(n \log n)$
- greedy-choice property:** let j^* be the item with *maximum* value/kg, v_j/w_j . Then there exists an optimal knapsack containing $\min(w_{j^*}, W)$ kg of item j^* .
- optimal substructure:** if we remove w kg of item j from the optimal knapsack, then the remaining load must be the optimal knapsack weighing at most $W - w$ kgs that one can take from $n - 1$ original items and $w_j - w$ kg of item j .

Proof. cut-and-paste argument

Suppose the remaining load after removing w kgs of item j was *not* the optimal knapsack weighing ... Then there is a knapsack of value $> X - v_j \cdot \frac{w}{w_j}$ with weight ... Combining this knapsack with w kg of item j gives a knapsack of value $> X \Rightarrow$ contradiction!

Minimum Spanning Trees

for a connected, undirected graph $G = (V, E)$, find a spanning tree T that connects all vertices with minimum weight. Weight of spanning tree T , $w(T) = \sum_{(u,v) \in T} w(u, v)$.

- optimal substructure:** let T be a MST. remove any edge $(u, v) \in T$. then T is partitioned into T_1, T_2 which are MSTs of $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$.

Proof. cut-and-paste: $w(T) = w(u, v) + w(T_1) + w(T_2)$ if $w(T') < w(T_1)$ for G_1 , then $T' = \{(u, v)\} \cup T'_1 \cup T_2$ would be a lower-weight spanning tree than T for G . \Rightarrow contradiction, T is the MST

- Prim's algorithm** - at each step, add the least-weight edge from the tree to some vertex outside the tree
- Kruskal's algorithm** - at each step, add the least-weight edge that does *not* cause a cycle to form


Binary Coding

Given an alphabet set $A : \{a_1, a_2, \dots, a_n\}$ and a text file F (sequence of alphabets), how many bits are needed to encode a text file with m characters?

- fixed length encoding:** $m \cdot \lceil \log_2 n \rceil$
 - encode each alphabet to unique binary string of length $\lceil \log_2 n \rceil$
 - total bits needed for m characters = $m \cdot \lceil \log_2 n \rceil$
- variable length encoding**
 - different characters occur with different frequency \Rightarrow use fewer bits for *more frequent* alphabets
 - average bit length, $ABL(\gamma) = \sum_{x \in A} f(x) \cdot |\gamma(x)|$
- BUT overlapping prefixes cause indistinguishable characters

Prefix coding


- a coding $\gamma(A)$ is a **prefix coding** if $\nexists x, y \in A$ such that $\gamma(x)$ is a prefix of $\gamma(y)$.
- labelled binary tree:** $\gamma(A)$ = label of path from root



- for each prefix code A of n alphabets, there exists a binary tree T on n leaves such that there is a **bijective mapping** between the alphabets and the leaves
- $ABL(\gamma) = \sum_{x \in A} f(x) \cdot |\gamma(x)| = \sum_{x \in A} f(x) \cdot |depth_T(x)|$
- the binary tree corresponding to an *optimal* prefix coding must be a **full binary tree**.
 - every internal node has degree exactly 2
 - multiple possible optimal trees - most optimal depends on alphabet frequencies
- accounting for alphabet **frequencies:**
 - let a_1, a_2, \dots, a_n be the alphabets of A in non-decreasing order of their frequencies.
 - a_1 must be a leaf node; a_2 can be a sibling of a_1 .
 - there exists an optimal prefix coding in which a_1 and a_2 are siblings
- derivation of optimal prefix coding: **Huffman's algorithm**
 - keep merging the two least frequent items

Huffman(C):

```
Q = new PriorityQueue(C)
while Q:
    allocate a new node z
    z.left = x = extractMin(Q)
    z.right = y = extractMin(Q)
    z.val = x.val + y.val
    Q.add(z)
return extractMin(Q) // root
```




10. REDUCTIONS & INTRACTABILITY

Reduction

Consider two problems A and B , A can be solved as follows:

- convert instance α of A to an instance of β in B
- solve β to obtain a solution
- based on the solution of β , obtain the solution of α .
- \Rightarrow then we say A **reduces** B .



instance \rightarrow another word for input

e.g. Matrix Multiplication & Squaring


- MAT-MULT: matrix multiplication
 - input:* two $N \times N$ matrices A and B .
 - output:* $A \times B$
- MAT-SQR: matrix squaring
 - input:* one $N \times N$ matrix C . *output:* $C \times C$
- MAT-SQR can be reduced to MAT-MULT
 - Proof.** Given input matrix C for MAT-SQR, let $A = C$ and $B = C$ be inputs for MAT-MULT. Then $AB = C^2$.
- MAT-MULT can also be reduced to MAT-SQR!
 - Proof.** let $C = \begin{bmatrix} 0 & A \\ B & 0 \end{bmatrix}$
 $\Rightarrow C^2 = \begin{bmatrix} 0 & A \\ B & 0 \end{bmatrix} \begin{bmatrix} 0 & A \\ B & 0 \end{bmatrix} = \begin{bmatrix} AB & 0 \\ 0 & BA \end{bmatrix}$

T-Sum

- o-SUM: given array A , output $i, j \in (1, n)$ such that $A[i] + A[j] = 0$
- T-SUM: given array B , output $i, j \in (1, n)$ such that $B[i] + B[j] = T$
- reduce T-Sum to o-Sum:**
 - given array B , define array A s.t. $A[i] = B[i] - T/2$.
 - if i, j satisfy $A[i] + A[j] = 0$, then $B[i] + B[j] = T$.

$p(n)$ -time Reduction

- $p(n)$ -time Reduction** \rightarrow if for any instance α of problem A of size n ,
 - an instance β for B can be constructed in $p(n)$ time
 - a solution to problem A for input α can be recovered from a solution to problem B for input β in time $p(n)$.
- ! n is in bits!**
- if there is a $p(n)$ -time reduction from problem A to B and a $T(n)$ -time algorithm to solve problem B , then there is a $T(O(p(n))) + O(p(n))$ time algorithm to solve A .



- $A \leq_p B \rightarrow$ if there is a $p(n)$ -time reduction from A to B for some polynomial function $p(n) = O(n^c)$ for some constant c . (" A is a special case of B ")

- if B has a polynomial time algorithm, then so does A
- "polynomial time" \approx reasonably efficient
- $A \leq_P B, B \leq_P C \Rightarrow A \leq_P C$

Polynomial Time

- **polynomial time** \rightarrow runtime is polynomial in the **length of the encoding** of the problem instance
- **"standard" encodings**
 - binary encoding of integers
 - list of parameters enclosed in braces (graphs/matrices)
- **pseudo-polynomial** algorithm \rightarrow runs in time polynomial in the **numeric value** if the input but is **exponential** in the **length** of the input
 - e.g. DP algo for KNAPSACK since W is in numeric value
- KNAPSACK is NOT polynomial time: $O(nW \log M)$ but W is not the number of bits
- FRACTIONAL KNAPSACK is polynomial time: $O(n \log n \log W \log M)$

Decision Problems


- **decision problem** \rightarrow a function that maps an instance space I to the solution set $\{YES, NO\}$
- decision vs optimisation problem:
 - **decision problem**: given a directed graph G , is there a path from vertex u to v of length $\leq k$?
 - **optimisation problem**: given ..., what is the *length* of the shortest path ... ?
 - convert from **decision** \rightarrow **optimisation**: given an instance of the optimisation problem and a number k , is there a solution with value $\leq k$?
- the decision problem is *no harder than* the optimisation problem.
 - given the optimal solution, check that it is $\leq k$.
 - if we cannot solve the decision problem quickly \Rightarrow then we cannot solve the optimisation problem quickly
- decision \leq_P optimisation

Reductions between Decision Problems

given two decision problems A and B , a polynomial-time reduction from A to B denoted $A \leq_P B$ is a

transformation from instances α of A and β of B such that

1. α is a YES-instance of $A \iff \beta$ is a YES-instance of B
2. the transformation takes polynomial time in the size of α



Examples

- **INDEPENDENT-SET**: given a graph $G = (V, E)$ and an integer k , is there a subset of $\leq k$ vertices such that no 2 are adjacent?
- **VERTEX-COVER**: given a graph $G = (V, E)$ and an integer k , is there a subset of $\leq k$ vertices such that each edge is incident to *at least one* vertex in this subset?
- **INDEPENDENT-SET \leq_P VERTEX-COVER**
 - **Reduction**: to check whether G has an independent set of size k , we check whether G has vertex cover of size $n - k$.

Proof. If INDEPENDENT-SET, then VERTEX-COVER.

Suppose (G, k) is a YES-instance of INDEP-SET. Then there is subset S of size $\geq k$ that is an independent set.

$V - S$ is a vertex cover of size $\leq n - k$. Proof: Let $(u, v) \in E$. Then $u \notin S$ or $v \notin S$.

So either u or v is in $V - S$, the vertex cover.

Proof. If VERTEX-COVER, then INDEPENDENT-SET.

Same as above, but flip IS and VC


e.g. SET-COVER

Given integers k and n , and collection S of subsets of $\{1, \dots, n\}$, are there $\leq k$ of these subsets whose union equals $\{1, \dots, n\}$?

Claim: **VERTEX-COVER \leq_P SET-COVER**


Reduction: given (G, k) instance of VERTEX-COVER, generate an instance (n, k', S) of SET-COVER.

Proof. For each node v in G , construct a set S_v containing all its outgoing edges. (Number each edge)



e.g. 3-SAT

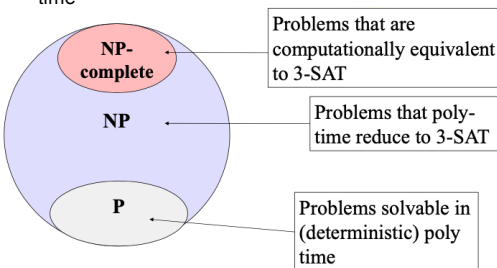
- **SAT**: given a CNF formula Φ , does it have a satisfying truth assignment?
 - literal: a boolean variable or its negation x, \bar{x}
 - clause: a disjunction (OR) of literals
 - conjunctive normal form (CNF): formula Φ that is a conjunction (AND) of clauses
- **3-SAT** \rightarrow SAT where each clause contains exactly 3 literals
- **3-SAT \leq_P INDEPENDENT-SET**
 - **Reduction**: Construct an instance (G, k) of INDEP-SET s.t. G has an independent set of size $k \iff \Phi$ is satisfiable
 - node: each literal term
 - edge: connect 3 literals in a clause in a triangle
 - edge: connect literal to all its negations
 - reduction runs in polynomial time
 - \Rightarrow for k clauses, connecting k vertices form an independent set in G .



11. NP-COMPLETENESS

- **P** \rightarrow the class of *decision* problems solvable in (deterministic) polynomial time
- **NP** \rightarrow the class of *decision* problems for which polynomial-time verifiable **certificates** of YES-instances exist.
 - aka *non-deterministic polynomial*

- i.e. no poly-time algo, but verification can be poly-time
- **certificate** \rightarrow result that can be checked in poly-time to verify correctness
- $P \subseteq NP$: any problem in **P** is in **NP**.
- if $P = NP$, then all these algos can be solved in poly time



NP-Hard and NP-Complete

- a problem A is said to be **NP-Hard** if for every problem $B \in NP, B \leq_P A$.
 - aka A is at least as hard as every problem in **NP**.
- a problem A is said to be **NP-Complete** if it is in **NP** and is also **NP-Hard**
 - aka the hardest problems in NP.
- **Cook-Levin Theorem** \rightarrow every problem in NP-Hard can be poly-time *reduced* to 3-SAT. Hence, **3-SAT is NP-Hard and NP-Complete**.
- NP-Complete problems can still be approximated in poly-time! (e.g. greedy algorithm gives a 2-approximation for VERTEX-COVER)

showing NP-Completeness

1. show that X is in NP. \Rightarrow a YES-instance has a certificate that can be verified in polynomial time
2. show that X is NP-hard
 - by giving a poly-time reduction from another NP-hard problem A to X . $\Rightarrow X$ is at least as hard as A
 - reduction should *not* depend on whether the instance of A is a YES- or NO-instance
3. show that the reduction is valid
 - 3.1. reduction runs in poly time
 - 3.2. if the instance of A is a YES-instance, then the instance of X is also a YES-instance
 - 3.3. if the instance of A is a NO-instance, then the instance of X is also a NO-instance

```
def INDEPENDENT-SET(G, k) -> bool:
1. G', k' = reduction(G, k)
2. yes_or_no: bool = CLIQUE(G', k') # magically given
3. return yes_or_no
```

What to show for a **correct** reduction:

- (G, k) is YES-instance $\rightarrow (G', k')$ is also a YES-instance
- (G', k') is YES-instance $\rightarrow (G, k)$ is also a YES-instance
- The transformation takes polynomial time in the size of (G, k)

showing NP-HARD

1. take any **NP-Complete** problem A
2. show that $A \leq_P X$

helpful approximations

stirling's approximation: $T(n) = \sum_{i=0}^n \log(n-i) = \log \prod_{i=0}^n (n-i) = \Theta(n \log n)$

harmonic number, $H_n = \sum_{k=1}^n \frac{1}{k} = \Theta(\lg n)$

basel problem: $\sum_{n=1}^N \frac{1}{n^2} \leq 2 - \frac{1}{N} \xrightarrow{N \rightarrow \infty} 2$
because $\sum_{n=1}^N \frac{1}{N^2} \leq 1 + \sum_{x=2}^{\log_3 n} \frac{1}{(x-1)x} = 1 + \sum_{n=2}^N (\frac{1}{n-1} - \frac{1}{n}) = 1 + 1 - \frac{1}{N} = 2 - \frac{1}{N}$
number of primes in range $\{1, \dots, K\}$ is $> \frac{K}{\ln K}$

asymptotic bounds

$1 < \log n < \sqrt{n} < n < n \log n < n^2 < n^3 < 2^n < 2^{2n}$
 $\log_a n < n^a < a^n < n! < n^n$
for any $a, b > 0$, $\log_a n < n^b$

multiple parameters

for two functions $f(m, n)$ and $g(m, n)$, we say that $f(m, n) = O(g(m, n))$ if there exists constants c, m_0, n_0 such that $0 \leq f(m, n) \leq c \cdot g(m, n)$ for all $m \geq m_0$ or $n \geq n_0$.

set notation

$O(g(n))$ is actually a *set of functions*. $f(n) = O(g(n))$ means $f(n) \in O(g(n))$
• $O(g(n)) = \{f(n) : \exists c, n_0 > 0 \mid \forall n \geq n_0, 0 \leq f(n) \leq cg(n)\}$
• $\Omega(g(n)) = \{f(n) : \exists c, n_0 > 0 \mid \forall n \geq n_0, 0 \leq cg(n) \leq f(n)\}$
• $\Theta(g(n)) = \{f(n) : \exists c_1, c_2, n_0 > 0 \mid \forall n \geq n_0, 0 \leq c_1 \cdot g(n) \leq f(n) \leq c_2 \cdot g(n)\} = O(g(n)) \cap \Omega(g(n))$
• $o(g(n)) = \{f(n) : \forall c > 0, \exists n_0 > 0 \mid \forall n \geq n_0, 0 \leq f(n) < cg(n)\}$
• $\omega(g(n)) = \{f(n) : \forall c > 0, \exists n_0 > 0 \mid \forall n \geq n_0, 0 \leq cg(n) < f(n)\}$

example proofs

Proof. that $2n^2 = O(n^3)$
let $f(n) = 2n^2$. then $f(n) = 2n^2 \leq n^3$ when $n \geq 2$.
set $c = 1$ and $n_0 = 2$.
we have $f(n) = 2n^2 \leq c \cdot n^3$ for $n \geq n_0$.

Proof. $n = o(n^2)$
For any $c > 0$, use $n_0 = 2/c$.

Proof. $n^2 - n = \omega(n)$
For any $c > 0$, use $n_0 = 2(c + 1)$.

Example. let $f(n) = n$ and $g(n) = n^{1+\sin(n)}$.
Because of the oscillating behaviour of the sine function, there is no n_0 for which f dominates g or vice versa.
Hence, we cannot compare f and g using asymptotic notation.

Example. let $f(n) = n$ and $g(n) = n(2 + \sin(n))$.
Since $\frac{1}{3}g(n) \leq f(n) \leq g(n)$ for all $n \geq 0$, then $f(n) = \Theta(g(n))$. (note that limit rules will not work here)

mentioned algorithms

- ch.3 - **Misra Gries** - space-efficient computation of the majority bit in array A
- ch.3 - **Euclidean** - efficient computation of GCD of two integers
- ch.3 - **Tower of Hanoi** - $T(n) = 2^n - 1$
 - move the top $n - 1$ discs from the first to the second peg using the third as temporary storage.
 - move the biggest disc directly to the empty third peg.
 - move the $n - 1$ discs from the second peg to the third using the first peg for temporary storage.
- ch.3 - **MergeSort** - $T(n) = T(\lfloor n/2 \rfloor) + T(\lceil n/2 \rceil) + \Theta(n)$
- ch.3 - **Karatsuba Multiplication** - multiply two n -digit numbers x and y in $O(n^{\log_2 3})$
 - worst-case runtime: $T(n) = 3T(\lceil n/2 \rceil) + \Theta(n)$

uncommon notations

- \perp - false

NP-Complete problems

1. Circuit-SAT
2. CNF-SAT
3. 3-SAT
4. Independent Set
5. Vertex Cover
6. Max-Clique
7. Hamiltonian Cycle (both directed and undirected)
8. Hamiltonian Path (undirected)
9. Traveling Sales Person
10. Subset Sum
11. Knapsack
12. Hitting Set