

Lectures

L1: Introduction

Four Vs of Data Science

- Volume
- Variety
- Velocity
- Veracity - uncertainty of data


Storage Hierarchy

- Volume: Server *le* Rack *le* Cluster
- Speed: Server *ge* Rack *ge* Cluster

Bandwidth vs Latency

- Throughput** Actual rate at which data is transmitted across the network over a period of time
- Bandwidth** Maximum (capacity) amount of data that can be transmitted per unit time
- Latency** Time taken for 1 data packet to go from source to destination (or both ways)
- Latency does not matter when transmitting a large amount of data
- Bandwidth does not matter when transmitting a small amount of data

Cost of moving data



- Bandwidth drops and latency increases as we move up the data hierarchy
- Disk reads are also much more expensive

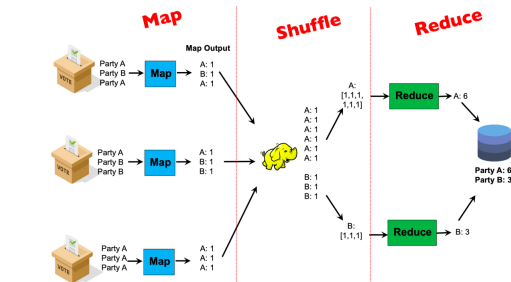
Big ideas of data processing

- Horizontal scaling is cheaper than vertical scaling
- Move data processing to the machine with the data since data clusters have limited bandwidth
- Process data sequentially and avoid random access to reduce total seek time
- Seamless scalability → use more machines to reduce time taken to process data

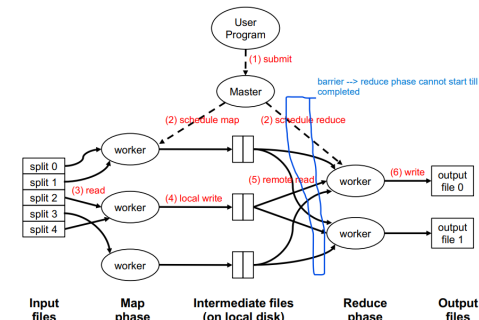
Challenges

- Machine failures
- Synchronisation
- Programming difficulty

L2 Map reduce



- Map: extract something of interest from each. Emits a key value pair
- Shuffle: Shuffle intermediate results by key value pairs
- Reduce: Aggregate intermediate results
- Each of these three processes can occur concurrently across different machines



Map Reduce Implementation

- Submit:** User submits mapreduce program and configuration (e.g. no. of workers) to Master node
- Schedule:** Master schedules resource for map and reduce tasks (master does not handle actual data)
- Read:** Input files are separated into splits of 128MB. Each split corresponds to one map task. Each worker executes map tasks 1 at a time
- Map phase:** Each worker iterates over each key,value tuple and applies the map function
- Local write:** each worker writes the output of map to intermediate files on its local disk. These files are partitioned by key
- Remote read:** each reduce worker is responsible for ≥ key. For each key, it reads the data it needs from the corresponding partition of each mapper's local disk
- Write:** output of the reduce function is written (usually to a distributed file system such as HDFS)

Interface

- map(k,v) → list(k',v')
- reduce

L3: No SQL Overview

NoSQL

- Not Only SQL: can include sql
- Stores data in a format other than relational DB
- Sql refers to relational DBMS, not the querying language - NoSQL can have querying lang too

- Used for large volumes of data and data that does not fit in a structured data (e.g. some has image, some don't)

Properties

- Horizontal Scalable: easy to partition and distribute across machines
- Replicate and distributed over many servers
- Simple call interface
- Often weaker concurrency model than RDBMS
- Efficient use of distributed indexes and RAM
- Flexible schema

Major NoSQL DB

- Key-stores
 - Stores mapping (associations) between keys and values
 - Keys are usually primitives (int, str, raw bytes etc) that can be easily queried
 - Values can be primitive or complex; usually cannot be easily queried (lists, JSON, HTML, BLOB)
- Operations
 - Get - fetch value with key
 - Put - set value with key
 - Multi-Get, multi-put, range queries (must be comparable, e.g. int, str)


Suitable for

- Small continuous read and writes
- Storing basic information or no clear schema
- Complex queries are rarely required
- E.g. Storing user sessions, caches, user data that is often processed individually
- Implementation
 - Non-persistent: Just a big in memory hash table (E.g. redis, memcached) that needs to be regularly backed up to disk
 - Persistent: data is stored persistently to disk (E.g. RocksDB, Dynamo, Riak)

- Wide-column databases - stored sparsely

	Column family 1		Column family 2	
	Column 1	Column 2	Column 1	Column 2
Row key 1				
Row key 2				

- Rows describe entities
- Related groups of columns are grouped as column families (similar to separate tables, except they share the same row)
- Sparsity: If a column is not used for a row, it doesn't use space (saves space for sparse data)
- Document stores



- no schema (flexible schema)
- A data base can have multiple collections
- A collection (tables) can have multiple documents (rows)
- A document is a JSON-like object with field (columns) and values

- Different documents can have different field and can be nested
- Querying**
 - Unlike key val stores, doc stores allow querying based on the content
 - If the field does not exist on the doc, we just skip it when doing CRUD


- Graph databases

- Need to store information about the nodes and edges

- Vector Databases

- Store vectors (each row is a point in d dimensions)
- Usually dense, numerical, and high-dimensional
- Allow fast similarity search via locality sensitive hashing (LSH), similar to min-hashing
- Scalable, real-time updates, replication
- Good for LLM and vision models as they need to be converted to vectors, and search, recommendation, clustering can be easily added
- E.g. Milvus, Radis, MongoDB, Atlas, Weaviate

Consistency



- Strong: Any reads on all observers immediately read the same result after update (uses locks, higher latency)
- Eventual: If the system is working and we wait long enough, eventually all reads will produce the same value (correctness affected)


BASE

- Basically Available - basic writing and reading operations are available most of the time
- Soft state: without guarantees, we only have some probability of knowing the state at any time
- Eventually consistent
- Eventual consistency offers better availability at the cost of weaker consistency
- NoSQL allows for weaker consistency guarantees, and can be tuned to be stronger (tunable consistency)
- Suitable for statistical queries and social media feed but not suited for financial transactions

Duplication

- Motivation: Support join statements → how do we join 2 tables to form 1 new table
- Some optimizations in SQL may not be possible in NoSQL
- Denormalization:**
 - Storage is cheap! Duplicate data to boost efficiency
 - Tables are designed around potential join queries (pre-create the join tables)
 - Good if the queries types are fixed
 - What if a field is updated? → changes need to be propagated to multiple table


Pros



+ **Flexible / dynamic schema:**
suitable for less well-structured data

+ **Horizontal scalability:** we will discuss this more next week

+ **High performance and availability:** due to their relaxed consistency model and fast reads / writes



Cons

- **No declarative query language:** query logic (e.g. joins) may have to be handled on the application side, which can add additional programming

- **Weaker consistency guarantees:** application may receive stale data that may need to be handled on the application side

- Depends on:
 1. if denormalization is suitable
 2. importance of consistency
 3. complexity of queries (joins Vs read/write)